An Introduction to Using the QCDSP Computer

Bálint Joó and Robert D. Mawhinney

July 4, 2002

Contents

1	Introduction				
	1.1	Organisation	4		
	1.2	Notational conventions	4		
2	Hell	o World	5		
	2.1	Foo or Bar?	5		
	2.2	Front End and Logging In	5		
	2.3	QOS and Environment Variables	7		
	2.4	The QC-Shell and Booting the QCDSP	9		
	2.5	Booting the QCDSP	10		
	2.6	Hello World	15		
	2.7	Summary	18		
3	Basi	c Programming of the QCDSP	20		
	3.1	Introduction	20		
	3.2	Parallelism in the QCDSP Computer	20		
	3.3	Brief description of the QCDSP Architecture	21		
	3.4	Programming Interface	23		
	3.5	Who Am I?	25		
	3.6	Nearest Neighbour Communications	28		
	3.7	Getting To Know Your Neighbours	34		
	3.8	Summary of Chapter	37		
	3.9	Programming Exercise	39		

4	Coll	ective Communication	41
	4.1	What are Collective Communications?	41
	4.2	Who Communicates?	41
	4.3	Global Communication in General	42
	4.4	Global Communication on the QCDSP	46
	4.5	Global Reduction Algorithms for Ring Architectures	47
	4.6	A simple collective communications library	53
	4.7	Summary of Chapter	57
	4.8	Exercise	57
5	I/O	on the QCDSP	58
	5.1	Introduction	58
	5.2	Parallel I/O	58
	5.3	Parallel I/O with NTFs, gread and gload	59
	5.4	Creating Node Tagged files using gread	66
	5.5	Loading data with qload	70
	5.6	Digression for Lattice Folk	75
	5.7	Serial I/O	86
	5.8	Summary	86
6	Con	apiling, Linking and Libraries	91
	6.1	Overview	91
	6.2	Tartan C++ Software Tools	92
	6.3	The C++ Compiler	92
	6.4	QCDSP Runtime and IO Libraries	95
	6.5	Converting TOFF to COFF	95
	6.6	A typical Makefile	95
	6.7	Linker Control Files	96
	6.8	The Tartan Archiver	100
	6.9	Summary	100

7	Single Processor Optimisation I – Using CRAM			
	7.1	Introduction	102	
	7.2	Overlay Programming and the Linker Control File	103	
	7.3	Step One: Straight C++	108	
	7.4	Step 2: Relocate C++ to CRAM	112	
	7.5	Step 3: Assembler optimisation	114	
	7.6	Summary	121	

Chapter 1

Introduction

These notes are an attempt at gently introducing a new user to the QCDSP computer. It reflects my own particular path of learning (which may or may not do me credit).

1.1 Organisation

We begin in chapter 2 where we discuss some introductory aspects of the QCDSP, including how to boot a QCDSP machine and how to run a simple "Hello World" application.

1.2 Notational conventions

Generally, when illustrating dialogue with the computer we shall use the teletype font. Filenames shall usually be denoted in **bold face**.

Chapter 2

Hello World

2.1 Foo or Bar?

It is the long and honourable tradition of user guides to begin by a walk through example of how to get a simple program (usually named **foo.c** or **bar.c** in more reputable texts) to run on a certain development platform. Such a program usually does something simple like write the message "Hello World" to the screen. In this chapter we should like to show how to get such a program running on a QCDSP computer.

2.2 Front End and Logging In

The QCDSP is a parallel system that is accessible through a *front end* host, which is a possibly simpler, serial computer such as a Sun Workstation.

This is an accepted principle in supercomputing circles. For example, the Thinking Machines Corporation Connection machines were often accessed from a Sun front end, as were some of the Elenia–Quadrics supercomputers. Cray T3E computers often have another Cray supercomputer (such as a J90 or a YMP) as their front end.

It is important to point out, however, that the Connection Machine and Quadrics machines were true *back end* machines in the sense that user programs ran on the front end Sun workstation, and the parallel hardware was controlled entirely by the front end, usually for array operations. The Cray and QCDSP supercomputers are different. They compute independently of the front end. The front end however provides important services such as a compilation environment, control to boot the parallel machines and disk space for the data produced by the QCDSP.

The QCDSP front end discussed in the rest of this section is a Sun workstation running the SunOS 4 variant of the UNIX operating system. Before being able to log in one requires an account to be set up on the system, which should be requested from one of the systems staff of the QCDSP

collaboration¹.

Login Procedures

In the following example I shall use a host called qcdhost.phys.columbia.edu, which is the front end to two QCDSP boards (each containing 64 processors). The QCDSP boards are themselves known as $q_{-}1$ and $q_{-}2$. This machine is also the front end to some other machines of which more will be said later.

Logging in to QCDSP is done according to the standard network login procedure. One can use telnet to reach the host from anywhere on the Internet although nowadays the more secure SSH login is recommended, which apart from forwarding X-Windows connections also encrypts the network communications and does strict host—name resolutions. As an example I use my current machine which is named *lattice2* to log into *qcdhost* via ssh.

```
[bj@lattice2 bj]$ ssh bj@qcdhost.phys.columbia.edu
bj@qcdhost.phys.columbia.edu's password:
Last login: Wed Mar 29 11:15:05 2000 from lattice2.pa.uky.
SunOS Release 4.1.4 (QCDHOST.sz) #2: Fri Nov 20 14:32:57 EST 1998
(qcdhost:/homeqs0/bj)%
```

Note that the password needs to be entered at the prompt in the second line, but, as is characteristic of UNIX the password is not echoed back to the terminal. If ssh is not available one can use telnet:

```
[bj@lattice2 bj]$ telnet qcdhost.phys.columbia.edu
Trying 128.59.196.11...
Connected to qcdhost.phys.columbia.edu.
Escape character is '^]'.

SunOS UNIX (qcdhost)

login: bj
password:
Last login: Wed Mar 29 11:20:05 2000 from lattice2.pa.uky.
SunOS Release 4.1.4 (QCDHOST.sz) #2: Fri Nov 20 14:32:57 EST 1998 (qcdhost:/homeqs0/bj)%
orrlogin:
```

¹Currently Bob Mawhinney, email rdm@phys.columbia.edu

```
[bj@lattice2 bj]$ rlogin -lbj qcdhost.phys.columbia.edu
bj@qcdhost.phys.columbia.edu's password:
Last login: Wed Mar 29 11:25:05 2000 from lattice2.pa.uky.
SunOS Release 4.1.4 (QCDHOST.sz) #2: Fri Nov 20 14:32:57 EST 1998
```

although in the latter two cases the protocols send the user passwords over the networks unencrypted and liable to interception by pesky hacksters worldwide. We note that when using *rlogin* one ends up talking to a *SSH* connection at the QCDSP end.

2.3 QOS and Environment Variables

QOS Version

The QCDSP computer runs an operating system entitled *QOS*. Numerous revisions of *QOS* exist with the current version at the time of writing this note being *QOS* 5.3.3. A UNIX shell environment variable, QOS_VERS is used to identify the version of *QOS* that a given user will be using. This environment variable should be set automatically be the shell startup files at the time the user logs in. In particular *qcdhost* supports the *TC-Shell* (tcsh) and the startup file in which the environment variables are set is **.cshrc**.

The toplevel QOS directory on qcdhost is

```
/qcdsp/sfw/qos.<version>,
```

where version refers to the operating system revision held by the variable QOS_VERS. For version 5.3.3 the relevant directory is

```
/qcdsp/sfw/qos.5.3.3.
```

The Tartan C++ Compiler

Another set of environmental variables that need to be set are related to the installation of the Tartan C++ compiler. The two variables in my .cshrc file at the time of writing are TC_DIR and TC_BOARD. Currently these two variables are set to the following values:

```
TC_DIR=/qcdsp/sfw/qos.<version>/usr:/usr/local/tartan/v2.1
TC_BOARD=qcdsp_v<version>
```

where version is the value of the variable QOS_VERS. I believe that TC_DIR is a list of directories which contain subdirectories for #include files, run time libraries etc. I suspect that TC_BOARD identifies, which QOS operating system version interfaces the hardware is compliant with (but this is guesswork on my part. Bob, please enlighten us here.) In any case these variables need to be set to sensible values for code to compile, link, resolve and run successfully on the QCDSP computer.

The QOS Toplevel Directory

Let us now pass a cursory glance at the toplevel QOS directory. A simple listing of files therein reveals the following directory contents

```
(qcdhost:/qcdsp/sfw/qos.5.3.3)% ls
Makefile
                                  make.log
README
                                  os 03 20 00 18:36.tar.qz
                                  plaq/
announce
buq/
                                  aio/
diag/
                                  qio_qker/
distrib/
                                  gker/
example/
                                  rts/
include/
                                  usr/
lib/
                                  util_prog/
machines/
```

Currently for us the most interesting of these are going to be the subdirectories called machines and usr. In the former are defined the various machine partitions available, whereas in the latter reside some useful header files describing the parallel programming communications objects.

Looking at the machines directory reveals the following list:

```
(qcdhost:/qcdsp/sfw/qos.5.3.3)% cd machines/
(qcdhost:/qcdsp/sfw/qos.5.3.3/machines)% ls
Makefile
               cc_1/
                              q4_{1}
                                              q6_2/
                                                              q_1/
c4 1/
               cc_2/
                              q4_{2}
                                              q7_1/
                                                              q_2/
c4 2/
               cc 3/
                              q5_{1}
                                              q7_2/
                                                              src/
cb_2/
               q3_{1}
                              q5_{2}
                                              q8_1/
cb_3/
               q3_2/
                                              q8_{2}
                              q6_1/
```

Each directory with the exception of src refers to a QCDSP processor partition. We are qoing to use machines q_1 and q_2 for the rest of this guide so let us examine q_1 :

```
(qcdhost:/qcdsp/sfw/qos.5.3.3/machines/q_1)% ls
boot_n0_notest.qrb
                        notest.qrb
                                                 q_1_sec4/
boot_n0_notest_25.qrb
                        notest_25.qrb
                                                 q_1_sec5/
                                                 q_1_sec6/
debug.qrb
                        q_1@
debug 25.grb
                        q_1_pri0/
                                                 standard.grb
debug_notest.qrb
                        q_1_sec1/
                                                 standard_25.qrb
debug_notest_25.qrb
                        q_1_sec2/
dev
                        q_1_sec3/
```

Most of the files in this directory have the suffix .qrb. These files are the configuration files for booting the QCDSP computer. They specify which operating systems kernels are to be sent to

the QCDSP daughterboards, what speed the inter processor communications should run at, what self test procedures to carry out at boot time, *etc*. In normal usage, the file **standard.qrb** should probably be used (the manner of its use will be described in the next section).

The file **q_1** is a symbolic link to a program called the *QC-Shell*, which is the shell environment used to boot, and run programs on the QCDSP, of which we shall say more below.

2.4 The QC-Shell and Booting the QCDSP

The QC-Shell - qcsh

So far we have been happily romping around the filesystem of the front end host. Now comes the time to actually boot a set of processors for our use.

Interaction with the QCDSP generally takes place through the aforementioned *QC-Shell*. This effectively is the usual *TC-Shell* with some bolted on features (in the form of built in commands) to manipulate the QCDSP computer. Commands for manipulating the QCDSP begin with the letter q. Some useful commands which we shall be using later are <code>qreset_boot</code>, <code>qrun</code> and <code>qprintf</code>.

We note the distinction between a shell built in command and UNIX commands. For example in a normal UNIX *TC-Shell* the command echo is a built in command of the shell, while the usual UNIX command 1s is an executable program usually residing in the directory /bin. Hence the command 1s is available to all shells, whereas the command echo is shell specific (although it is implemented in almost all UNIX shells).

Like most UNIX–shell like shells the *QC-Shell* allows for self initialisation through a startup file. In the case of the *QC-Shell* this file is called **.qcshrc**. If your account has been set up properly you should find a **.qcshrc** file in your toplevel homedirectory. My one reads as follows:

```
source ~/.cshrc
qinit
set curdir = 'pwd'
source $curdir/dev
qset_qd_verbose 2
qset_qd_scsi_tree_verbose 2
qset_reset_boot f standard.qrb
qset_exit_info 5 exit.info
qreset_boot
```

One of the chief functions of the .qcshrc file is to boot the QCDSP computer. This is accomplished by various commands in the .qcshrc file.

The qinit command initialises the QCDSP system and starts up the Qdaemon system daemon which is responsible for for booting the QCDSP machine, loading and executing programs, per-

forming self tests and other important tasks. By "sourcing" the **dev** file in the current directory, several variables are defined in the shell specifying amongst other things the boot–debug network device (in this case SCSI ID) which is used to communicate with the QCDSP boards. **Note that this .qcshrc expects you to be in the correct machine directory at the time of invocation of the QC-Shell, where the file dev resides.**

Thereafter some boot time verbosity levels are set to specify how much of the various Qdaemon messages and other test results should be sent to the standard output for boot time diagnostic purposes. These levels are set by the qset_qd_verbose and the qset_qd_scsi_tree_verbose commands.

The configuration (.qrb) file to use when booting, or resetting the QCDSP is specified with the command <code>qset_reset_boot</code>. In my case it is the **standard.qrb** file which was mentioned in the last section.

The qset_exit_info command itself is used to what's it used to do then?

Once all these settings have been put into effect the machine is finally booted through the command qreset_boot. Incidentally this is also a useful bail out command. Should the QCDSP crash or deadlock, or for whatever reason we may want to reboot the computer, it can be done using the qreset_boot command in the *QC-Shell*.

More information on these commands will be given in the advances stages of this manual (if it ever gets that far...)

2.5 Booting the QCDSP

Now that we have examined how one boots the QCDSP, let us do it by invoking the *QC-Shell*. The magic incantation for machine $q_{-}1$ is:

```
(qcdhost:/homeqs0/bj)% cd /qcdsp/sfw/qos.5.3.3/machines/q_1
(qcdhost:/qcdsp/sfw/qos.5.3.3/machines/q_1)% ./q_1
```

The first command takes to the machine directory, where the **dev** and **.qrb** files live. Recalling that q_1 is a symbolic link the qcsh command, it can be seen that the second command invokes the *QC-Shell*.

Aside: Since getting to the correct machine directory involves a lot of typing it might be useful to set up some sort of shell alias to allow us reaching the directory more quickly. In my **~.cshrc** file I have two variables set:

```
set q_1 = /qcdsp/sfw/qos."$QOS_VERS"/machines/q_1
set q_2 = /qcdsp/sfw/qos."$QOS_VERS"/machines/q_2
```

Hence I can reach the q_1 and q_2 directories respectively by just typing cd q_1 and cd q_2 respectively.

The invoking of the *QC-Shell* should start a large set of messages scrolling up your screen, giving you the results of boot time diagnostics and self tests. Some important ones to look out for:

```
qinit: qcsh initialized

Qdaemon state is:
    qdaemon task is INITIALIZE_DAEMON
    qdaemon abort no and resume no
    QCDSP not synchronous
    SCSI tree coordinates selected: mb 0, db 0,
    Load address 0 (hex), data 0 (hex) from screen
    Read address (hex) 0, blksize 0 (hex) from screen
    Read output to screen in node tagged format
    Run entry address from file
    SCSI packet and payload size: 512 488
    Run kernels not installed on all nodes
    Hardware debugging level is 0
    Not decoding packet headers
```

This message signifies that the *QC-Shell* has been initialized and the so called Qdaemon has been started. The Qdaemon is the primary source of interaction with the QCDSP. Note also the Qdaemon message has been displayed twice, once for entry into the specific service routine and once at the exit. This is the result of setting the verbosity level in the **.qcshrc** file. Carrying on with the boot process the following messages should appear at some point:

These messages show the results of trying to talk via the SCSI bus to the QCDSP, in particular that the SCSI bus has been found and that the Qdaemon has now entered the machine reset and boot service. If all is well the QCDSP should reply as:

Szm::inquiry: COLUMBIA QCDSP
ResetBoot::ResetBoot: Mb 0 found
ResetBoot: Working on scsi layer 0
ResetBoot: turning off LED's on layer 0

Devel Development in the control of the control of

ResetBoot: starting non SCU NGA setup on layer 0

ResetBoot: starting DramStdTest0 on layer 0
ResetBoot: DramStdTest0 passed for layer 0
ResetBoot: loading run kernels to layer 0

If you do not get the COLUMBIA QCDSP message then something is wrong. The rest of the messages display various test results.

The various motherboards are arranged in a tree structure with boot time communication being done via a SCSI bus. One of the next tasks of the boot process is to find out about the SCSI tree. This information is also echoed back at boot time (given the verbosity levels in my startup files):

Summary of SCSITree object SCSITree: SCSIPath information for motherboard 0 SCSI ID path: SCSI branch path: SCSIBranch information for branch 0 SCSIBranch: information for branch 0 For SCSI ID 0 the device is motherboard 0 (serial number 38) at slot 0 and crate 0xf1 For SCSI ID 1 no device is present For SCSI ID 2 no device is present For SCSI ID 3 no device is present For SCSI ID 4 no device is present For SCSI ID 5 no device is present For SCSI ID 6 no device is present For SCSI ID 7 no device is present SCSILayer information for layer 0

SCSILayer: Motherboards in this layer

Hence we see that in booting $q_{-}1$ there is only one motherboard on SCSI Id 0.

The machine then proceeds to boot all the daughter boards in turn. This is done by sending boot kernels to all the daughter boards. Useful messages to look out for are

DSPSerialInit: getting information for all motherboards

DSPSerialInit: All MB's responded with no errors

DBBoot: Booting all Db's on all Mb's

All nodes 1 to 63 on MB 0 booted and responded

DBBoot: All MB's responded after Db boot

ResetBoot: finished send of boot kernels to dbs

DSPSerialInit: getting information for all motherboards

DSPSerialInit: All MB's responded with no errors

DSPSerSwitchTest: Starting to check status for all motherboards

Network switched over all routes 10 times.

Maximum of 10 reads done to clear receive buffer

Mb 0 passed test

DSPSerSwitchTest: All requested MB's responded.

All DSP serial connections OK.

that show that all the motherboards and daughterboards booted OK and that the serial connections have been checked and are responding. If the machine gets this far, the boot kernels boot the node gate arrays (NGA-s) and perform a DRAM memory test. If this succeeds the runtime kernels are downloaded to the daughterboards as evidenced by messages such as:

ResetBoot: starting NGA boot for dbs

ResetBoot: finished NGA boot for dbs

ResetBoot: starting DRAM stdtest for dbs

Checking DRAM stdtest for dbs on Mb 0

ResetBoot: DRAM Stdtest on all dbs passed

ResetBoot: loading run kernels to dbs

Finally the machine is synchronised and the Serial Communications are set up. At this point some communications tests are performed (even and odd wire tests). Healthy messages look along the lines of

ResetBoot::SCUSetup: Beginning SCU setup of all nodes

ResetBoot::Sync: Beginning to synchronize QCDSP

ResetBoot::Sync: Checking that all nodes have sync set

ResetBoot::Sync: QCDSP synchronous

ResetBoot: running even wire send SCU test

ResetBoot: running odd wire send SCU test

ResetBoot: SCU test passed

ResetBoot::InitNodeStat: starting send of NodeStat
ResetBoot::InitNodeStat: starting build of NodeTable

NodeTable::NodeTable: constructor invoked for 1 motherboards

ResetBoot::InitNodeStat: NodeTable complete

Finally if all goes well the Qdaemon will report that it has exited the boot service by the message

```
Qdaemon state is:

qdaemon task is R
```

qdaemon task is RESET_BOOT
qdaemon abort no and resume no
QCDSP synchronous
SCSI tree coordinates selected: mb 0, db 0,

and the QC-Shell will likewise show its status message

```
qcsh after RESET_BOOT:
All nodes selected with SCSI tree coordinates
```

and the **boot process is complete**. Your shell prompt should turn into something along the lines of

```
(qcdhost/qcdsp/sfw/qos.5.3.3/machines/q_1: qcsh[q_1])%
```

Congratulations! You have now booted a QCDSP supercomputer. Some things to be aware of

- The *QC-Shell* does not really like lots of processes going on in parallel. Recall that it is designed to interact with the QCDSP. This shell really should be run only to submit QCDSP jobs. Compiling, editing, e-mailing and the like should be carried out from the usual login *TC-Shell*.
- While in the *QC-Shell* you are blocking everyone else's access to a potential HPC resource. (From a 3.2Gflop board up to a Tflops scale supercomputer.) Spend your time in it wisely.
- If things become messy you can always type <code>qreset_boot</code> to reboot the system.

2.6 Hello World

Compiling the Program

The time has now come to execute the *Hello World* application. The easiest method is to take a copy from the directory of programming examples. You can do this by typing:

```
(qcdhost:/homeqs0/bj)% cd /qcdsp/sfw/qos.5.3.3/example/
(qcdhost:/qcdsp/sfw/qos.5.3.3/example)% tar cf - hello_world | ( cd ; tar x x hello_world/Makefile, 608 bytes, 2 tape blocks
x hello_world/main.C, 1233 bytes, 3 tape blocks
x hello_world/hello_world.ctl, 188 bytes, 1 tape blocks
x hello_world/main.tof, 3076 bytes, 7 tape blocks
x hello_world/main.lst, 2113 bytes, 5 tape blocks
x hello_world/hello_world.map, 14453 bytes, 29 tape blocks
x hello_world/hello_world.outtof, 15276 bytes, 30 tape blocks
x hello_world/hello_world.out, 14654 bytes, 29 tape blocks
```

This should copy the **hello_world** directory to your homedirectory. A nice makefile has been provided so that we can build the application by executing the commands:

```
(qcdhost:/homeqs0/bj)% cd hello_world/
to go to the source directory and then:
```

```
(gcdhost:/homegs0/bj/hello_world)% make clean
```

to remove the currently built stuff so we can do a clean rebuild. The computer should respond with

```
rm -f *.tof *.lst *.map core *.qin *.out *.outtof *.ctl \ .
```

We perform the build by using the command make:

```
(qcdhost:/homeqs0/bj/hello_world)% make
```

to which the computer should reply:

```
tcpp -c -mb -km -kc -kl -q main.C
tcpp -e hello_world.outtof -mb -km -kc -kl -q main.tof \
    tcio30bs.olb
Tartan Linker for QCDSP version 5.3.2, modified 1/11/99 by RDM
Memory = 0x80000, stack = 0x1000, heap = 0x1000
```

```
(See transcript written on "/homeqs0/bj/hello_world/hello_world.map" for additional information.)
t2c hello_world.outtof hello_world.out
```

In the first line, the make system instructs the Tartan C++ tcpp compiler to compile the source file **main.C** which is the source for the actual *hello_world* application. The second invocation of tcpp links the result of the compilation with the run time system to produce a file in the common object format. Finally the t2c command creates an executable format suitable for the QCDSP. The QCDSP executable is the one called **hello_world.out**.

Running The Application

Once the application has been built, it should be run. First boot the QCDSP. This time around we will use machine q_2 .

after booting has completed we get the prompt:

```
(qcdhost/qcdsp/sfw/qos.5.3.3/machines/q_2: qcsh[q_2])%
```

To run the hello_world application we use the qrun command:

```
(qcdhost/qcdsp/sfw/qos.5.3.3/machines/q_2: qcsh[q_2])% cd ~/hello_world/
(qcdhost/homeqs0/bj/hello_world: qcsh[q_2])% qrun hello_world.out
```

The results should be a bunch of Qdaemon messages, indicating that the executable is being loaded to the QCDSP, followed by the output from Motherboard 0, daughterboard 0:

```
Hello world

Motherboard and daughterboard numbers for this node: 0 0

Unique ID for this node: 0

Physics coordinates for this node: 0, 0, 0, 0

The square of 0 is 0

The square of 1 is 1

The square of 2 is 4

The square of 3 is 9

The square of 4 is 16

The square of 5 is 25

The square of 6 is 36
```

```
The square of 7 is 49
The square of 8 is 64
The square of 9 is 81
```

followed by various exit information messages.

The output of node 0 (Motherboard 0, Daughterboard 0) is always echoed to the standard output. The output of other nodes are saved to various buffers on the nodes. We can retrieve the saved output from the processors using the <code>qprintf</code> command, which sends the saved messages to the standard output. After some Qdaemon messages the results should look like

```
Odaemon state is:
        qdaemon task is PRINTF
        qdaemon abort no and resume no
        QCDSP synchronous
        All nodes selected with SCSI tree coordinates
Qdaemon: user print buffer for Mb 0 and Db 0 (SCSI tree coordinates)
Buffer wrap count 0
Hello world
Motherboard and daughterboard numbers for this node: 0 0
Unique ID for this node:
Physics coordinates for this node: 0, 0, 0, 0
The square of 0 is 0
The square of 1 is 1
The square of 2 is 4
The square of 3 is 9
The square of 4 is 16
The square of 5 is 25
The square of 6 is 36
The square of 7 is 49
The square of 8 is 64
The square of 9 is 81
         user print buffer for Mb 0 and Db 1 (SCSI tree coordinates)
Odaemon:
Buffer wrap count 0
Hello world
Motherboard and daughterboard numbers for this node: 0 1
Unique ID for this node:
                          1
Physics coordinates for this node: 0, 0, 0, 1
The square of 0 is 0
The square of 1 is 1
The square of 2 is 4
```

```
The square of 3 is 9
The square of 4 is 16
The square of 5 is 25
The square of 6 is 36
The square of 7 is 49
The square of 8 is 64
The square of 9 is 81
```

and so on for all the processors on the machine.

Congratulations! If you have followed all that I detailed here you have now wasted the computer time of a 3.2GFlop peak speed development board to print "Hello World" and some squares of numbers on the screen.

Exiting QCDSP

To exit fromt the QCDSP simply logout of the *QC-Shell* by typing exit.

2.7 Summary

After discussing logging in to the QCDSP system. I said a few introductory remarks about QOS and the environmental variables needed in the shell to be able to use QCDSP properly. I then described the boot process for the machine. Finally I have gone through the process of building the example *Hello_World* application, running it and retrieving the standard output from the various QCDSP processor buffers to where they were saved.

This is the minimal subset of knowledge needed to be able to write and execute programs on the QCDSP computer.

QOS and Shell Environment Summary

The login startup files should set the environmental variables QOS_VERS, TC_DIR and TC_BOARD. The current QOS version at the time of writing this document is version 5.3.3.

The toplevel QOS directory is /qcdsp/sfw/qos.<version>. The available machine partitions are in /qcdsp/sfw/qos.<version>/machines. The QCDSP specific include files are in /qcdsp/sfw/qos.<version>/usr/include and the QCDSP run time libraries are in /qcdsp/sfw/qos.<version>/usr/lib, where version is the value held in the QOS_VERS.

Booting Summary

Go to the relevant machine directory, and invoke the *QC-Shell* using the symbolic link provided (usually the same as the machine name). To use machine q_2 for example the incantation is

```
(qcdhost:/homeqs0/bj)% cd /qcdsp/sfw/qos.5.3.3/machines/q_2
(qcdhost:/homeqs0/bj)% ./q_2
```

The machine should then boot if all goes well.

Compilation Summary

Currently we did not detail compilation. We simply used the makefiles provided, which set the required compilation flags, and performed the linking. We recommend copying these makefiles for now.

The compiler is named topp, which is the Tartan C++ compiler. There is also a utility t2c which turns the output of the compilation and linking into an executable suitable for the QCDSP. The QCDSP executables produced by the makefile have the **.out** suffix.

Running Summary

Once the executable has been built, it can be run using the qrun command. The output from processor 0 (Motherboard 0, Daughterboard 0), is automatically echoed to the standard output. The saved standard output buffers from the QCDSP processors can be retrieved to the standard output of the front end host using the qprintf command.

Miscalleny

The QCDSP computer can be reset with the <code>qreset_boot</code>. The QCDSP machine can be freed up for others by simply logging out of the *QC_Shell* using the <code>exit</code> command.

This concludes the "Hello World" chapter. In the next chapter we shall discuss interprocessor communications.

Chapter 3

Basic Programming of the QCDSP

3.1 Introduction

This chapter is intended as a guide to the basic system calls implemented on the QCDSP computer.

The QCDSP itself is most simply programmed in C or C++, for which a commercial compiler exists with some optimisation features. The compiler is called the Tartan C++ compiler and the UNIX invocation of it is via the command tcpp.

The machine can also be programmed at the assembler level and indeed such programming may be necessary to achieve maximal computational and communication performance. However, for the moment we shall not discuss this mode of programming.

3.2 Parallelism in the QCDSP Computer

The QCDSP, whether at the level of a single development board or at the level of several cabinets containing processor boards is a MIMD¹ parallel computer.

This means, that each processor is capable of running completely independent code. However, the machine is perhaps best programmed according to the SPMD² programming model. In this programming model processors run the same program, using different datasets resident to each processor. Also this programming model allows one to write fully MIMD programs, by executing separate subroutines (of the same program) depending on some identification token (a unique processor ID, processor grid coordinates or some such) from the local processor.

Furthermore, the QCDSP has a distributed memory, with no current implementation of a shared memory system. The preferred method of communication is via message passing between neighbouring processors. We shall outline possible methods to implement global communications such

¹Multiple Instruction Multiple Data

²Single Program Multiple Data

as broadcasts and global sums in the next chapter. For the current chapter, we shall introduce some of the QCDSP system calls to allow a given particular processor to identify itself, and to send messages to its neighbours. To make the discussion clearer, we have to mention a few details about the architecture and network topology of the QCDSP computer.

3.3 Brief description of the QCDSP Architecture

The QCDSP at the bottom level consists of a collection of Texas Instruments TMS320C31 Digital Signal Processors (DSPs) with associated DRAM memory (0.5Gwords of error correcting memory per processor) and special purpose communications hardware. These components are organised into a hierarchy of *daughterboards*, *motherboards* and *crates*.

A daughterboard consists of a single DSP chip, its associated DRAM and its so called Node Gate Array (NGA), which is the communications unit built on a custom chip. A single daughterboard is about the size of a credit card. A picture of a daughterboard is shown in figure 3.1

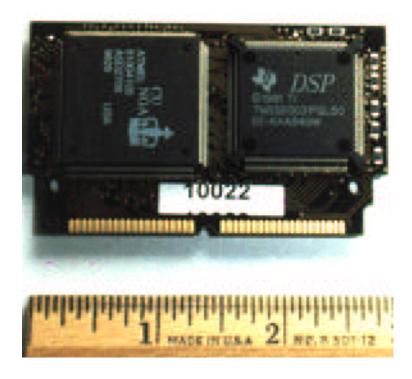


Figure 3.1: Photograph of a daughterboard. The DSP chip and the NGA are shown with a ruler to indicate size in inches

The daughterboards are themselves mounted on motherboards. A single motherboard holds up to 64 daughterboards. Of these, one daughterboard (daughterboard 0) is integrated onto the motherboard and has some special functions that we shall describe later. A picture of a motherboard can be seen in figure 3.2.

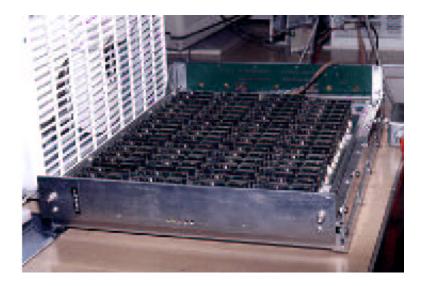


Figure 3.2: Photograph of a motherboard containing 63 daughterboard cards and one integrated daughterboard. The motherboard is mounted on a test rig and is cooled by the fan on the left hand side

The daughterboards are connected on the motherboard in a four dimensional (4D) mesh via a *serial network* which is intended to be the main medium of communications for applications and hence it has been named *the physics network*. The physics network is driven by the NGA of each daughterboard. This 4D serial network wraps around in each dimension, giving each motherboard network topology of a *4D*–*torus*.

All the daughterboards on a given motherboard are also connected to the special daughterboard 0 (the one integrated on the motherboard) via a second serial network in a star topology. This network is part of the so called *boot–diagnostic network*. The boot–diagnostic network is used for hardware testing, and for downloading the various boot and runtime kernels to the daughterboards at boot time.

Each motherboard also has two SCSI ports which are connected to daughterboard 0 on that board. These ports allow the connection of daughterboards to either a front end host, to another motherboard, or in future designs of the machine, to local disk systems. The SCSI connections form part of the boot–diagnostic network and are not in general expected to be heavily used in physics applications.

The motherboards are organised into crates. The *physics networks* of the boards in a given cabinet are connected by a tagliatelle tangle of flat ribbon serial cables into an overall physics network which still forms a 4D torus. The boot diagnostic networks of the boards in a crate are connected via the SCSI interfaces on the motherboards into a tree hierarchy. A picture of a crate is shown in figure 3.3.

Finally the crates are assembled into a full machine.

In summary the communications in a machine are made up of two networks:



Figure 3.3: A QCDSP crate holding 6 motherboards

- The Physics Network which is a serial network that to the programmers point of view forms a 4D torus. This is the network that is intended for primary communications in an application.
- The Boot–Diagnostic network made up from a SCSI chain between motherboards, rooted on the front end workstation. On each motherboard the boot–diagnostic network is constructed from a serial network with a star topology rooted on daughterboard 0.

A summary of the networks in QCDSP can be seen in figure 3.4.

3.4 Programming Interface

The QCDSP at the simplest level provides a set of C++ functions which may be called to achieve various aims. Broadly these functions can be categorised as

• Functions allowing the processor to identify itself and the machine

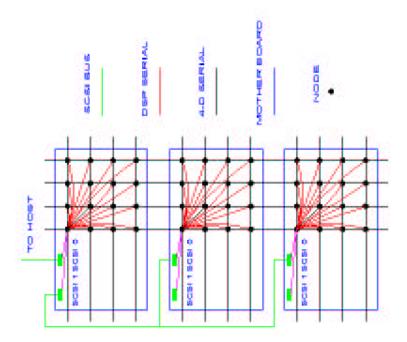


Figure 3.4: A summary diagram of QCDSP Networks showing three motherboards connected in a tree of depth = 2. The horizontal and vertical black lines show the physics network (2 dimensions of it) with the black dots showing the daughterboards. The green lines show the SCSI tree and the red lines show the serial boot–diagnostic network. Purple lines show the connections between the SCSI ports and motherboard 0 on each board.

- Objects and functions for carrying out communications
- functions for self synchronisation
- miscellany random numbers etc.

Header Files

These functions are defined in the header file **sysfunc.h** in the directory /qcdsp/sfw/qos.<version>/usr, where <version> refers to the QOS revision as discussed in the previous chapter. This directory should be listed as part of the TC_DIR shell environment variable (see last chapter), thus making the directory part of the default search path for the Tartan C++ compiler. The programmer merely needs to include the header file in his source files using the directive

#include <sysfunc.h>

identically to the way he/she normally includes **stdio.h**, **stdlib.h**, **math.h** and other standard header files.

Unsupported C++ Features

One should also be aware that not all of the standard ANSI C++ libraries are implemented on the QCDSP. In particular the iostream classes are unavailable. The programmer needs to use the C standard I/O routines defined in **stdio.h**.

Compilation and Running of Examples

In the following sections we describe some of the functions provided in the **sysfunc.h** file. To run the example codes we recommend that the reader copies the *hello_world* program as described in the last section into his homedirectory and modifies the file **main.C**. The makefile provided can then be used to build the examples by typing the command make. The resulting executable will be called **hello_world.out**. (Readers with a knowledge of makefiles can adapt the makefile as they wish. The compiler flags and command line options for the C++ compiler can be listed by typing the command tapp with no arguments.) Let us assume from here on that the *hello_world* directory was copied to one of identical name in the users homedirectory and that the results of all compilations are named **hello_world.out**.

3.5 Who Am I?

Knowing your identity and place in the Universe

The following functions are defined in **sysfunc.h** to enable a processor to localise itself:

```
int UniqueID() :
```

This function returns an int specifying the **unique processor ID** of the calling processor. (N.B: The MPI equivalent of this call would be MPI_Comm_Rank).

```
int NumNodes() :
```

This function returns an int specifying the total number of nodes that are in use in the current machine partition. (N.B: The MPI equivalent of this call would be MPI_Comm_Size).

```
int MbNum() :
```

This function returns an int specifying the **motherboard number** of the motherboard on which the calling processor resides.

```
int DbNum() :
```

This function returns an int specifying the **daughterboard number** of the calling processor.

```
int CoorT(), int CoorX(), int CoorY(), int CoorZ() :
```

These function calls return int-s specifying the T, X, Y and Z coordinates respectively, of the calling processor within the 4D torus of the physical network.

```
int SizeT(), int SizeX(), int SizeY(), int SizeZ() :
   These functions return int-s specifying the size of the 4D torus physical network in the T,
   X, Y and Z directions respectively
```

Figure 3.5 shows a programming example, where the processors in the QCDSP identify themselves and size up their machine parition using these functions:

```
#include <stdio.h>
#include <stdlib.h>
#include <sysfunc,h>
int main(int argc, char *argv[])
   /* ----- What is my Unique ID -----*/
   int my_id = UniqueID();
   int num_processors = NumNodes();
   /* -----*/
   int my motherboard num = MbNum();
   int my_daughterboard_num = DbNum();
      - My processor grid coordinates and grid dimensions -\star/
   int my_coords[4];
   int proc_grid_dims[4];
   /* ----- Get my coordinates ----- */
   my_coords[0] = CoorT();
my_coords[1] = CoorX();
   my_coords[2] = CoorY();
   my_coords[3] = CoorZ();
   /* - Get information about the processor grid size --- */
   proc_grid_dims[0] = SizeT();
proc_grid_dims[1] = SizeX();
   proc grid dims[2] = SizeY();
   proc grid dims[3] = SizeZ();
   /* - Echo Back information to the user ----- */
   printf("I am processor: %d.\n", my id);
   printf("I live on motherboard %d, daughterboard %d\n", my_motherboard_num,
                                                                   my_daughterboard_num);
   printf("I am one of a total of %d computing elements\n", num_processors);
   printf("The current physical network has the following dimensions \verb|\n"|);\\
  printf("%d Processors in the T direction\n", proc_grid_dims[0]);
printf("%d Processors in the X direction\n", proc_grid_dims[1]);
printf("%d Processors in the Y direction\n", proc_grid_dims[2]);
printf("%d Processorz in the Z direction\n", proc_grid_dims[3]);
   \verb|my_coords[0]|, \verb|my_coords[1]|, \verb|my_coords[2]|, \verb|my_coords[3]||;
   return(EXIT_SUCCESS);
```

Figure 3.5: Sample program showing use of self identification functions.

Running the code in figure 3.5 on the 64 daughterboard development board $q \cdot 2$ produced the following results:

```
2 Processorz in the Z direction My (T,X,Y,Z) coordinates are (0, 0, 0, 0 ) . .
```

Using the command aprintf command as described in the last chapter allows one to look at the output buffers of some of the other processors:

```
Qdaemon: user print buffer for Mb 0 and Db 2 (SCSI tree coordinates)
Buffer wrap count 0
I am processor: 2.
I live on motherboard 0, daughterboard 2
I am one of a total of 64 computing elements
The current physical network has the following dimensions
4 Processors in the T direction
4 Processors in the X direction
2 Processors in the Y direction
2 Processorz in the Z direction
My (T,X,Y,Z) coordinates are (0, 0, 1, 0)
Qdaemon: user print buffer for Mb 0 and Db 63 (SCSI tree coordinates)
Buffer wrap count 0
I am processor: 63.
I live on motherboard 0, daughterboard 63
I am one of a total of 64 computing elements
The current physical network has the following dimensions
4 Processors in the T direction
4 Processors in the X direction
2 Processors in the Y direction
2 Processorz in the Z direction
My (T,X,Y,Z) coordinates are (3, 3, 1, 1)
```

Some useful enumerated types

You may notice³ if you look through the code in 3.5 that I have always ordered the indices of both the my_coords and proc_grid_dims in the order T, X, Y and Z with the T corresponding to the index 0 and Z corresponding to index 3. To make for much more readable programs there exists an enumerated type called SCUAxis. A variable of type SCUAxis can take on only four values

```
SCU_X -- for the T direction,

SCU_X -- for the X direction,

SCU_Y -- for the Y direction and

SCU_Z -- for the Z direction.
```

The actual definition of the type:

³if you haven't yet please do so now

```
enum SCUAxis { SCU_T, SCU_X, SCU_Y, SCU_Z };
```

ensures that the integer values of the enumerations are SCU_T=0, SCU_X=1, SCU_Y=2 and SCU_Z=3 respectively.

Thus we should be able to index our arrays with these enumerations rather than with the integers 0,1,2 and 3 as before. For example we could write

```
my_coords[SCU_X]=CoorX(); /* equivalent to my_coords[1] = CoorX() */
```

or for example we could cycle through all the directions in the processor grid with the following for loop:

```
int i;
for(i = SCU_T; i <= SCU_Z; i++) {
   /* do something exciting here */
}</pre>
```

Incidentally, we note that this latter use relies on the particular ordering of the axes defined in the enumerated type. In future releases of QOS, the ordering of the enumerated elements maybe different from the current release. If this happens the above for loop may not be portable. Hence, *caveat emptor*. The enumerated type itself is defined in /qcdsp/sfw/qos.<version>/include/scu_enum.h which is included in the **sysfunc.h** file. (Suggestion: perhaps a member of the enumerated type called MAX_AXIS or something should be defined for such loops. They could then be written as for (i=0; i< MAX_AXIS; i++)

3.6 Nearest Neighbour Communications

The QCDSP **sysfunc.h** header files defines several ways of carrying out nearest neighbour communications. We will look into the simplest one here. By communication we mean sending/receiving a certain number of bytes from one processor to another. On the QCDSP such a communication is made up of three stages.

Preparation: – the sender/receiver describes the data to be sent/received and provides a pointer in memory to the data that is to be sent / to where the data is to be received

Initiation: – the sender / receiver initiates the communication. The sender / receiver is then free to carry out other tasks.

Completion: – the sender / receiver wait until the communication is complete, in the case of the sender this means that all the data has been sent and in the case of the receiver it means that the data has arrived.

Communication along the physical network is handled by the Serial Communication Unit (SCU) in the NGA chip, and the communication does not need much interaction from the CPU past the

initiation stage. Hence the time spent communicating allows the processor to carry out computation while the communication proceeds. It is only immediately before using the results of the communications that the processor has to wait to make sure that the communication is complete. (Note: This is similar in style to the MPI non–blocking communications, however in the case of MPI the so called preparation and initiation stages are usually rolled into one)

Preparation

In preparing to communicate we must decide upon the following things:

• Who did we want to communicate with? This is generally one of our neighbours. We have neighbours in eight directions, these being the positive and negative T, X, Y and Z directions respectively. An enumerated type, SCUDir, is made available to help us in writing safe, maintainable programs. This enumerated type is made accessible through the header file **sysfunc.h**. The actual enumerations are

```
SCU_TM (T, Minus): — Communicate in the negative T direction
SCU_TP (T, Plus): — Communicate in the positive T direction
SCU_XM (X, Minus): — Communicate in the negative X direction
SCU_XP (X, Plus): — Communicate in the positive X direction
SCU_YM (Y, Minus): — Communicate in the negative Y direction
SCU_YP (Y, Plus): — Communicate in the positive Y direction
SCU_ZM (Z, Minus): — Communicate in the negative Z direction
SCU_ZP (Z, Plus): — Communicate in the positive Z direction
```

• We must decide whether we wish to send a message or receive one. Once again we have an enumerated type at our disposal called SCUXR. This can have only two values:

```
SCU_REC: receive operation SCU_SEND: send operation.
```

• We must describe our data. For our purposes the data shall consist of a number of blocks, each block having a certain length. The blocks need not be contiguous in memory, but may have some regular spacing between them. The distance in memory (in units of blocks) separating the start of one block from the start of the next one is called the data stride.

Hence our messages are completely described by specifying the *memory address* of the first block, *the length of a block in bytes, the number of blocks in the message* and *the stride*. For messages made up of a contiguous set of blocks, the stride is always one. In other words, the start of a block in memory is always one block-length away from the start of the previous block with the intervening memory between the two being filled up by the body of the first block. The idea of this data description scheme is illustrated in figure 3.6.

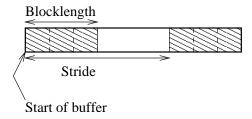


Figure 3.6: Anatomy of a message. Here the message consists of 2 blocks (shaded) with each block having a block length of 3 bytes. The blocks are not contiguous but have a block separation between them making the stride value for this message 2

On the QCDSP the preparing to communicate involves the instantiation of a so called SCUDirArg object. There are several ways of carrying out this instantiation. Perhaps the simplest is to use the *constructor*:

This method allows the complete description of the putative communication to be specified with one single statement. The arguments have the following meaning:

void *buffer this is a pointer to the start of the message in memory. It is of type void * so that the pointer can point to data of any type.

SCUDir dir this is the direction in which the communication is to proceed and should take the value of one of the 8 enumerations of the enumerated type SCUDir listed above.

SCUXR xr indicates whether we are sending or receiving. This should take one of the two values of the enumerated type SCUXR listed above.

int blklen gives the length in bytes of a single block to be transmitted/received.

int numblk = 1 gives the number of blocks to be communicated. This is an optional argument to the constructor, which, if unspecified will have the default value of 1, indicating that there is only a single block to be communicated

int stride = 1 specifies the stride of the data. This is an optional argument to the constructor, which, if unspecified will have the default value of 1 corresponding to the contiguous data

As an example, suppose I wanted to send an integer in the positive X direction, I could instantiate the communication description object as:

or as:

where in the last two cases I have explicitly specified the optional arguments. In all the above cases the result of the instantiation is an object of type SCUDirArg with name send_int_x_plus. This object now acts as a *handle* to that particular communication.

Likewise to receive a 5-vector of floating point numbers of type float from the negative Y direction, I would set up the communication as:

Alternatively, I can set up the communication to send the 5 members of a 10 member integer array which have odd indices in the positive T direction using the following instantiation:

The SCUDirArg objects can also be instantiated and manipulated using various other *class methods*. For example I can instantiate an object, without a description of the communication to be performed by simply using the default constructor:

```
SCUDirArg for_later;
```

Later on I can initialise this object with the method Init which has the same argument list as the full constructor described earlier. Hence I can either set up the object in the previous manner as:

or alternatively I can use the Init method:

Other methods by which the data description elements of the object can be set are

- void * Addr(void *addr) the start of the data is set to addr. The address of the new buffer is returned.
- int Blklen(int blklen) the block length of the communication is set to blklen. The new block length is returned.
- int Numblk (int numblk) the number of blocks in the transfer is set to numblk. The new number of blocks is returned.
- int Stride (int stride) the stride of the transfer is set to stride. The new stride is returned.
- void Reload(void *a, int blklen, int numblk=1, int stride = 1) The new data description indicated in the argument list is set (the send/receive nature of the data or the communication direction remain as before). a holds the address of the start of the data. The remaining arguments should be self explanatory by now. Note again that the last two arguments are optional.

So for example to send first an int, then a float in the positive X direction followed by a second int communication, I can instantiate the communication object for the first int and then simply use the methods above to change the data description:

```
/*
  * Carry out the second communication here
  */

/* Now prepare to send an int again -- last two arguments optional */
send_x_plus.Reload((void *)&int_buf, sizeof(int));

/*
  * Carry out the third communication here
  */
```

Initiating the Communications

Once we have defined the type of communication we wish to perform, we must initiate the data transfer. This is done by the function SCUTrans. There are several ways of using SCUTrans, here I will describe the simplest one.

Calling SCUTrans with the address of an SCUDirArg object as the sole argument will initiate the communication encoded in that object.

As an example, suppose we want to send the integer 4 in the positive X direction and an SCUDirArg object int_x_plus has been instantiated to describe this communication (*Exercise: Instantiate an SCUDirArg object to encode this communication*.) The communication could then be initiated by the following call:

```
SCUTrans( &int_x_plus );
```

However the results of the communication will not be available until the communication completes.

Completing the Communication

To complete the communication the program should call the function SCUTransComplete. One can call this function with the address of an SCUDirArg object to complete to communication encoded in that, or without any arguments to complete all outstanding communications. There is also another way to call SCUTransComplete to match a particular way of calling SCUTrans which will not be detailed here.

Hence to complete the previously initiated transfer encoded in the object x_int_plus of the last subsection, one should call

```
SCUTransComplete( &x_int_plus);
```

To complete all outstanding communications on a given processor (this does not imply a synchronisation across processors) one should instead call

3.7 Getting To Know Your Neighbours

Before concluding this chapter, we present the code for a complete program designed to show the use of some of the communication routines presented in the last section.

The program discovers the processor ID's of its neighbours in every direction. For each axis, T, X, Y and Z, the processors send their unique ID's in the positive direction, and receive the ID's of their neighbours from the negative direction.

Once this is done for all four axes, the communication pattern is repeated but in the opposite direction. That is, for all four axes, each processor will send its unique ID in the negative direction and receives the ID of its neighbour from the positive direction. When this process is done for all the axes, the processor will have the ID's of all its neighbours.

```
#include <stdio.h>
#include <stdlib.h>
#include <sysfunc.h>
// This function, when given the direction to send in (send direction)
// will send the contents of its send_buffer in that direction and
// will receive into its receive buffer from the negative direction
void neighbourSend(int *send_buffer, int *recv_buffer,
                  SCUDir send_direction)
   SCUDir recv_direction;
   // Figure out the receive direction
   switch( send_direction ) {
        case SCU_TM : recv_direction = SCU_TP;
                       break;
       case SCU_TP : recv_direction = SCU_TM;
                       break;
        case SCU_XM : recv_direction = SCU_XP;
                      break;
        case SCU_XP : recv_direction = SCU_XM;
                     break;
        case SCU_YM : recv_direction = SCU_YP;
                     break;
        case SCU_YP : recv_direction = SCU_YM;
                     break;
       case SCU_ZM : recv_direction = SCU_ZP;
       break;
```

```
case SCU_ZP : recv_direction = SCU_ZM;
    default: fprintf(stderr, "P%d: Error Invalid Send Direction\n");
             exit(EXIT_FAILURE);
             break;
 ,
// -----/
 // Set up SCUDirArg objects, one for send and one for receive //
 // ----- //
 // ------//
 // Send an integer in the send direction
 // ----- //
 SCUDirArg send( send_buffer,
          send_direction,
          SCU_SEND,
          sizeof(int));
 // ----- //
 // Recieve an integer from the opposite direction //
 // ----- //
 SCUDirArg recv( recv_buffer,
          recv_direction,
          SCU_REC,
          sizeof(int));
 // Initiate the transfer
 SCUTrans( &send );
 SCUTrans( &recv );
 // Wait for these two specific transfers to complete //
 // ------//
 SCUTransComplete( &send );
 SCUTransComplete( &recv );
 // ----- //
 // -----//
int main(int argc, char *argv[])
 // ----- //
 // Details about myself
 // ----- //
 int my_id=UniqueID();
 int my_coords[4] = { CoorT(), CoorX(), CoorY(), CoorZ() };
 // ----- //
 // Print out information about the grid
 // -----/
 printf("Processor grid has dimensions: T=%d X=%d Y=%d Z=%d\n",
     SizeT(),
     SizeX(),
     SizeY(),
     SizeZ());
 // ----- //
```

```
// Print out information about ourselves
printf("P%d, T=%d, X=%d, Y=%d, Z=%d\n",
     my_id,
     my_coords[SCU_T],
     my_coords[SCU_X],
     my_coords[SCU_Y],
     my_coords[SCU_Z]);
// ----- //
// Now Send my ID to my neighbour in each dimension //
// ----- //
int neighbour_id[8];
                     // -- Array to hold information about
                     //
                        neighbours
                     // A direction counter variable
SCUDir counter;
// ----- //
// This next bit is a little naughty as it relies on the fact //
// Ordering of the members of the enumerated types
// -----//
// Transmit my ID in the positive directions //
// and receive ID from the negative direction
                                               //
                                               //
// Note that Axis_plus = Axis_minus + 1 in the enumeration
                                               //
                                               //
// hence the choices of the counter
// First send our ID-s to the +ve directions (receive from -ve)
for(counter = SCU_TM; counter <= SCU_ZM; counter+=2) {</pre>
 neighbourSend( &my_id,
            &neighbour_id[ counter ],
            counter + 1 );
}
// Second send our ID-s to the -ve directions (receive from +ve)
// -----
for(counter = SCU_TP; counter <= SCU_ZP; counter+=2) {</pre>
 neighbourSend( &my_id,
            &neighbour_id[ counter ],
            counter - 1 );
}
// ----- //
// Print out my neighbours IDs
// -----//
printf("P%d: T-Neighbour (-,+): (%d, %d)\n",
     my_id,
     neighbour_id[ SCU_TM ],
     neighbour_id[ SCU_TP ] );
printf("P%d: X-Neighbour (-,+): (%d, %d)\n",
     my_id,
     neighbour_id[ SCU_XM ],
     neighbour_id[ SCU_XP ] );
```

3.8 Summary of Chapter

In this chapter we have outlined some of the basic system calls provided by the QCDSP **sysfunc.h** and how they can be used for processors to identify themselves and to perform simple nearest neighbour communications.

Architecture summary

The main point to stress about the computer architecture is that there are two main networks in the QCDSP. One is the boot–diagnostic network and the other is the physics network. The boot–diagnostic network is partly made up of a SCSI tree structure, partly of a serial network with a star topology rooted on the special embedded daughterboard on each motherboard. The physics network is a 4-D torus.

Self Identification Summary

Several routines are provided in the header file **sysfunc.h**. These include the functions

```
int UniqueID() - returns the processor ID
int NumNodes() - returns the number of processors
int MbNum() - returns the processors motherboard number
int DbNum() - returns the processors daughterboard number
int CoorT(), int CoorX(), int CoorY(), int CoorZ() - returns the processors coordinates in the physical network.
int SizeT(), int SizeX(), int SizeY(), int SizeZ() - returns the sizes of the physical network in the respective dimensions.
```

Enumerated Types Summary

The following enumerated types are defined in the hearder file sysfunc.h

```
enum SCUAxis = { SCU_X, SCU_X, SCU_Y, SCU_Z }; —a type to describe the four axes of the physical network
```

```
enum SCUDir = { SCU_TM, SCU_TP, SCU_XM, SCU_XP, SCU_YM, SCU_YP, SCU_ZM, SCU_ZP }; — a type to enumerate the 8 communication channels available on each processor. The naming convention is SCU_<Axis><P/M> where <Axis> is one of T, X, Y or Z and <P/M> is either P for the positive (plus) direction along the axis or M for the negative (minus) direction along the axis
```

enum SCUXR = { SCU_REC, SCU_SEND= 8}; — enumerates the two types of communications, sending or receiving. Note that SCU_SEND has an integer value of 8 rather than 1.

Nearest Neighbour Communications Summary

Nearest neighbour communications can be initiated by the SCUTrans function and completed by the SCUTransComplete function. The description is specified via an SCUDirArg object.

The SCUDirArg objects can be instantiated through invoking the constructor

where the arguments respectively are the address of the start of the data, the direction of the communication, the mode of communication (SEND/RECEIVE), the blocklength of the data to be communicated (in bytes), the number of blocks to communicate and the data stride. The last two of these parameters are optional. The default number of blocks is 1 indicating a single item transfer. The default stride is 1 indicating contiguous blocks.

Other forms of the SCUDirArg constructor are also available. (See more advanced documentation if I ever write it).

The SCUTrans function call has a prototype

```
void SCUTrans( SCUArgDir *arg );
```

where arg is a pointer to the SCUArgDir object describing the communication. A call to SCUTrans initiates a communication. The communication should not be considered complete until a corresponding SCUTransComplete call returns. The SCUTrans function has other (overloaded) prototypes which were not discussed here. (See more advanced documentation if I ever write it).

The SCUTransComplete function call has prototypes

```
void SCUTransComplete( SCUArgDir *arg );
void SCUTransComplete();
```

The first of these takes as its argument a pointer to an SCUArgDir object. In this case the function returns when the communication encoded in the object pointed to by arg completes. In the second case the function has no arguments. In this case the function returns when all the communications currently pending complete.

In the next chapter we shall discuss simple global communication algorithms and their implementation using the functions detailed in this chapter.

3.9 Programming Exercise

The Problem: – More efficient Neighbours

The program for discovering our neighbours can be made more efficient. Currently, a node only performs two communications at a time in each direction; one send and one receive. In principle, a processor could send its ID along all 8 available direction. Likewise, its neighbour could receive along all 8 of its communication channels. The goal of this exercise is for you to write a program that discovers its neighbours in this manner.

Checkerboard Partitioning

Checkerboard partitioning of processors (also known as even—odd partitioning or red—black partitioning) assigns a parity (colour) to each processor which can be either even (red) or odd (black). The name checkerboard partitioning comes from the fact neighbouring processors have opposite parities (colours). See figure 3.7 for a picture of a two dimensional processor grid separated into even and odd (red and black) sites.

The idea here that in the end processors of one colour (say red) can send their ID's to all their neighbours of the other colour (in this case black), while processors of the other colour (black) can receive along all 8 of their wires. Hence half of all the processors (the black ones) can discover the identities of all their neighbour in one single communication step. The communication could then be repeated, with the colours interchanged. In other words, in the second communication, all the processors of the second colour (black in this case) can transmit their IDs to processors of the first colour (red).

Thus all the nearest neighbours can be discovered using in effect two communication steps.

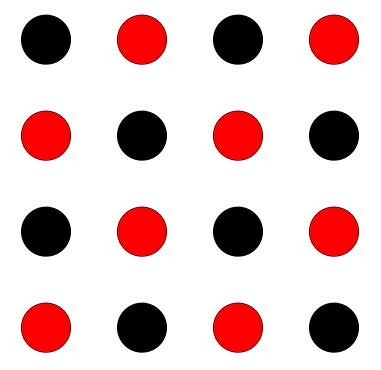


Figure 3.7: A 4x4 grid of processors with checkerboard partitioning. Note that any one processor has oppositely coloured neighbours

Exercise 1: Am I red or am I black

Let the two parities be denoted 0 (for red say) and 1 (for black). First thing the processor has to do is identify its parity. It could do this for example based on its coordinates in the physics network. Implement a routine using the relevant system calls by which each processor determines its parity.

Exercise 2: Do the transfer

Implement the neighbour discovery algorithm outlined above.

Chapter 4

Collective Communication

4.1 What are Collective Communications?

In the previous chapter we have briefly described the basic SCU calls for nearest neighbour point to point communication. This kind of communication is fine if we want to program communications between processors that are reminiscent of the HPC CSHIFT function for example, or if we wish to build systolic array like systems.

Occasionally however we may want to perform certain global communications such as a sum of data items held across all the processors. An immediate example of such a situation is the case of taking the scalar product of two vectors. Imagine that the vector components are distributed amongst the nodes of the parallel computer. Each processor could work out the sum of the squares of its local vector components but to accumulate the dot product, these local results would have to be globally summed.

4.2 Who Communicates?

The distinguishing feature of global communications also known sometimes as collective communications is that all the processors make the call to the communication routines, whereas in the case of a point to point communication, usually only the sending and the receiving node call a communication routine.

In the case of MPI and other message passing standards usually one has a way of selecting subgroups of processor nodes from a given available set. In MPI for example the so called 'communicator' MPI_COMM_WORLD refers to all available processors and subgroups of these can be identified by defining new communicators (using such routines as MPI_Comm_split.)

In the case of the QCDSP, selecting processor partitions is non trivial. Machines are often 'manually hardware partitioned' (a technical process involving graduate student plugging wires on the backplanes of crates). Alternatively, in a given hardware partition, one may select a certain set

of processors using the QOS command <code>qset_nodes</code>. Once a given set of nodes is selected by either method however, it is non-trivial to sub-select again from within a running program. Consequentially, our description of collective communications will refer to situations when all available processor nodes take part in the communication. These kinds of collective communications are referred to as global communications. From here on we shall use the two phrases interchangeably.

4.3 Global Communication in General

We now describe some common types of collective communication. Since a lot of these are in the MPI message passing interface specification I shall refer frequently to MPI equivalents. If you know MPI this should make things easy for you. It should also clear up ambiguities (I hope) in other situations.

Barrier

A barrier routine is used to synchronise a set of processors. All the nodes call the barrier routine. The routine does not exit until all the processes participating in the communication have entered it. In MPI one would require a call to the function MPI_Barrier. In the case of the QCDSP the functionality is provided through the system call sync(), which is defined in the system header file sysfunc.h.

Broadcast

There are two general kinds of broadcast routines. They are

One To All Broadcast: This is procedure where a given processor node (usually referred to as the root processor) contains some information. This information is sent to all other processors during a one to all broadcast routine.

All to All Broadcast: This is a procedure where at the outset each processor has an item of data, and during the broadcast procedure each processor makes available its data to all other processors. At the end of the broadcast, all the processors are in possession of the data elements from all the other processors.

The idea of the two kinds of broadcasts are illustrated in figures 4.1 and 4.2 respectively.

Important uses of broadcast operations can include the distribution of parameters from a source processor as well as situations where one processor has to be a 'master' and has to make global decisions. These decisions then probably need to be broadcast out to the other processors.

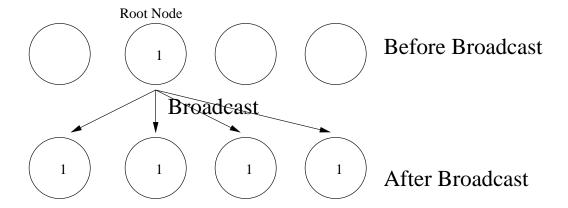


Figure 4.1: One to all broadcast. The circles represent processor nodes. Before the broadcast, the root node holds the data item to be distributed, in this case the number 1. After the broadcast, the data item is found on all processors

Gather and Scatter

The *gather* operation is one in which all processors send their data to a nominated root processor. Its opposite is the so called *scatter operation* where the nominated root processor distributes a vector of data, to the whole processor grid. The operations are illustrated in figure 4.3. Sometimes the terms *gather* and *scatter* are generalised, to mean that one processor distributes its data to some subset of other processors (for example its nearest neighbours in the processor grid). In this sense the corresponding gather would be the collection of data from the same subset of processors.

In fact the broadcast operation can be implemented as a scattering operation, where all processors receive the same data item. Likewise the neighbour discovery exercise at the end of the last exercise (have you attempted it?) is an example where we combine a generalised gather and a generalised scatter. The generalised gather is that each node (of a given parity) receives the processor ID of its nearest neighbours, whereas the generalised scatter is the part where every node (of the opposite parity) sends its ID to its nearest neighbours. The MILC collaboration, in their code¹ has refined the generalisation of gathering and scattering to such a level where a user can actually define his or her own gather and scatter mapping.

Scatter and Gather operations are useful for example to distribute data that has just been read from disk by a master processor (assuming of course that the master processor can hold all the data in memory) or to pass out processor specific decisions. It can also be used, as mentioned before to implement the broadcast operation. Generalised scatters can be useful for exchanging data boundaries between processors for example.

¹Freely available at http://cliodhna.cop.uop.edu/hetrick/milc/. You could do us a great service if you ported it to the QCDSP...

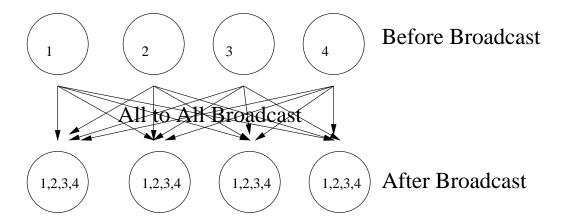


Figure 4.2: All to all broadcast: The circles represent processor nodes. Before the broadcast all items hold their own data (the numbers inside them). After the broadcast, each processor has a copy of everybody else's data.

Global Reduction

Global reduction operations generally take data from each node and produce one single (reduced) result. Examples of global reduction are global sums, global products, global XOR operations, finding global minima and maxima. In fact any associative operation can be used as a reducing operation. It perhaps for this reason that MPI has only two functions MPI_Reduce and MPI_Allreduce to carry out all their reduction operations², instead of having separate global sum, global XOR and other global routines.

Once again there are two kinds of global reduction operations.

All to One Reduction: – these are reduction operations where the final answer is left with one nominated (root) processor. These correspond to MPI_Reduce.

All to All Reduction: – these are reduction operations where the final answer is given to every processor. These correspond to MPI_Allreduce.

It should be clear for example that one can implement all to all reduction operations as an all to one reduction operation followed by a broadcast.

An important use of global sums for example is in the computation of scalar products. A use of a global maximum operation for example would be to find the infinity norm of a distributed vector. Global XOR operations may be needed for checksumming a distributed dataset. A quick and dirty floating point broadcast routine could be written using a global sum where the root node contributes to the sum the amount it wishes to broadcast and all the other processors contribute 0.

A warning about global sums and other global operations which are susceptible to cancellation errors and or bit overflow problems: The QCDSP system can contain anything from between 64

²The actual operation and type of data are specified using tags such as MPI_INT and MPI_SUM for example

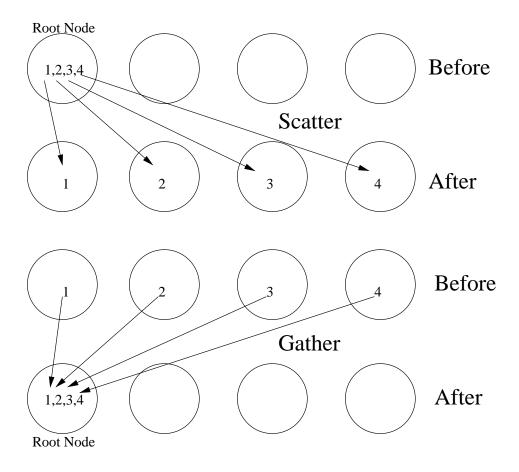


Figure 4.3: Above: Scatter operation. One processor 'scatters' its data amongst the complete set. Below: Gather operation: The data from all processors is 'gathered' onto a root processor.

nodes to $O(10^5)$ nodes. Even if a single node has only component to contribute to a particular reduction, it is possible that when for example 16K processors are involved in an operation there is a serious chance of overflow or that cancellation errors have a serious effect. There are two ways to address this problem:

- The implementors should implement their reduction operations in such a way as to minimise the problems of rounding error accumulation/ overflow/underflow. They could for example carry out sums in a logarithmic manner etc.
- The implementors could throw extra bits of precision at the problem, and provide status registers indicating overflow underflow.

4.4 Global Communication on the QCDSP

The QCDSP was designed to run simulations of lattice QCD. The predominant communication pattern for this application is nearest neighbour, with the occasional need for global operations. In fact lattice QCD does not really require more in terms of communications than broadcasts, nearest neighbour gather and scatter operations, global sums, global minima and maxima.

Consequentially, the design of the hardware and the software reflect these needs. The hardware design is such as to favour efficient nearest neighbour communications. Although the hardware does have support for performing global reduction operations, accessing them is not entirely straightforward. Also, since these global reduction operations were found not to be a major performance bottleneck, little effort has been made to optimise them. Even worse, the global communication routines that are optimised are part of the Columbia/Brookhaven Physics system and are not available for general use.

Distributed data is loaded to the processors by the QOS before applications start running. (More details in the Chapter on I/O) This sidesteps the need for routines that gather and scatter data to and from some root processor for the purposes of file I/O. Furthermore these are not entirely straightforward for a general user to code up as there is no hardware support for routing a message from any given processor to the root processor directly. An implementer would have to pass the message through intermediate processors, which are neither senders nor receivers of this information possibly having to stop computation on the intermediate processors for this purpose.

All that remains for a user to implement then are global reduction routines and broadcasts. We now look at some algorithms for performing these global communications. We then present a simple library that performs global sums, minima and maxima for numbers of type float. Users are welcome to use this library in their own programs, or can use the library merely as an example to aid in writing their own routines.



Figure 4.4: A 1D ring. The circles represent processors and the lines represent connections.

4.5 Global Reduction Algorithms for Ring Architectures

We now discuss some global reduction algorithms for 1D ring architectures. A 1D ring is simply a set of processors connected in series with wrap around edges at the ends of the line. A picture of a ring can be seen in figure 4.4.

Although the QCDSP has a network connectivity that is a four dimensional torus it can also be viewed as a collection of rings. We illustrate the idea for a two dimensional torus in figure 4.5. Here we can look at a two dimensional torus as a collection of rings in either the x or y directions

One general way of programming global reduction algorithms on a 4D mesh is to perform the operation first on all rings in the x direction, then on all the rings in the 'y' direction and so on in the remaining two directions.

Ring all to all reduction algorithm

Consider the following reduction algorithm:

Each processor in the ring sets up a buffer to store the result of the reduction. The result buffer is initialised with the processors own data.

Each ring then transmits its data in one direction along the ring and receives data from the opposite direction. (We shall refer to these as the positive and negative directions respectively). The data that has been received is combined with the data in the result buffer, to generate a new intermediate result..

The process of transmitting the local data in the positive direction and receiving data from the negative direction and combining it with the result is repeated until each data item has visited each processor once. For a ring containing n processors exactly n-1 steps are needed. When all the steps are complete, the result buffer should contain the final answer.

The process is illustrated for addition in a 4 processor ring in figure 4.6.

Note that in this particular algorithm, each processor performs the sum in a different order from its neighbour. This is one of the reasons why reduction operations need to be associative. Also since floating point arithmetic is not really associative due to roundoff errors, it is possible that rounding errors will affect different processors differently and for a floating point global sum using the above algorithm, it is possible that the answers on individual processors are not bit-identical.

This can be a problem for example if one is trying to decide whether one meets the stopping criteria for some iterative process based on the result of a global sum (such as in the case of an iterative

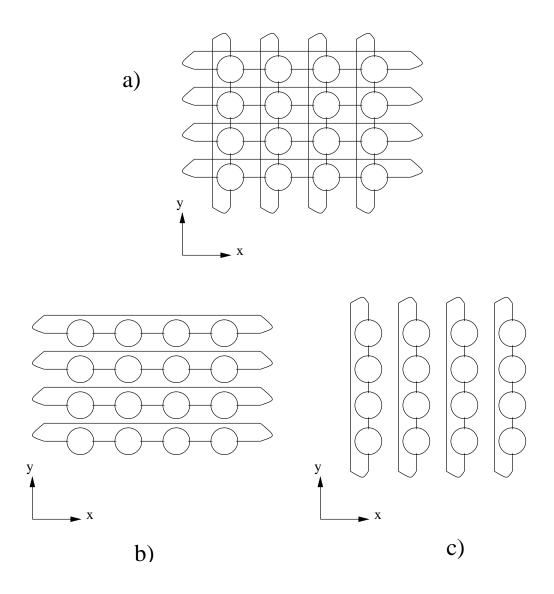


Figure 4.5: (a)A 2D torus as a set of 1D rings in the (b) x direction and (c) the direction

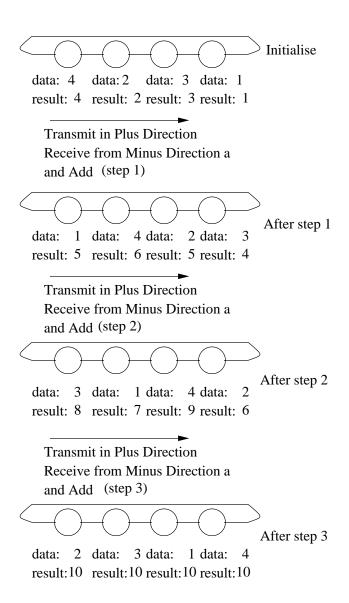


Figure 4.6: 1D Ring Addition Algorithm: At each step data is passed in the positive direction (right) and received from the negative direction (left). The data is then added to the results buffer. After 4-1=3 steps all the processors have the global sum



Figure 4.7: 2D Mesh Global sum. First the sum is performed along both rings in the horizontal direction. Second it is performed along both rings in the vertical direction.

solver, where the stopping criteria depends on a scalar product). Clearly it is possible that some processors meet this criteria and some, due to rounding errors do not.

The problem can be solved by nominating one processor as the master of the others. The decision as to whether to stop the solver will then depend on the result held by the master node. One can either broadcast the master's result and let each processor make its decision based on that, or alternatively the master can make the decision and broadcast it instead in the form of a token, a flag or in some other encoding. Alternatively one can formulate the global sum so that the results are guaranteed to be bit identical across all the processors, say by using a different algorithm where the sum is first reduced to some single root processor in all the rings which then transmit the sum back along their respective rings before going on to the next dimension.

Mesh Global reduction algorithm

As mentioned before, the generalisation of the ring algorithm to the mesh is simply to perform the global operation along each ring in parallel for a given direction, and to repeat this process for all directions.

Figure 4.7 shows the idea for a two dimensional mesh.

Coding the Routines For the QCDSP

We show below the QCDSP C++ code for finding the global maximum using similar the algorithms described above.

```
// -----
\#define \max(A, B) ((A) > (B) ? (A) : (B))
// -----
// Buffer space
// -----
static float transmit_buf;
static float receive_buf;
static float max_buf;
// -----
// Function glb_max
//
// Argument: A pointer to the local number that is to
// be considered for being the Global Maximum. At the end
// of the function, the pointer points to the global maximum
// -----
void glb_max(float * float_p)
{
 // -----
 // Size of the Processor Grid in each dimension
 // -----
 int NP[4] = {SizeT(), SizeX(), SizeY(), SizeZ()};
 // -----
 // Array to hold SCU send and receive directions
 // We will later index into this.
 // -----
 const SCUDir dir[] = { SCU_TP, SCU_TM, SCU_XP, SCU_XM,
                 SCU_YP, SCU_YM, SCU_ZP, SCU_ZM };
 // -----
 // Place the local number in the comparison buffer
 // -----
 max_buf = *float_p;
 // -----
 // Loop over the processor grid dimensions
 // -----
 int dim;
 for(dim = 0; dim < 4; dim++) {
    // -----
    // Send our local minimum in the +ve direction
    // (NP[ 2 * dim ]) direction and receive from the
```

```
// -ve (NP[ 2 * dim ]) direction with wraparound
  // at the pe grid boundary. Do this NP - 1 times
  // so that everyone can compare everybody's data
  // in this dimension
  // -----
  transmit_buf = max_buf;
  int tmp;
  for (tmp = 1; tmp < NP[ dim ]; tmp++) {
    // -----
    // Set up the communications handles:
    // -----
    SCUDirArg send(&transmit_buf, dir[ 2*dim ],
              SCU_SEND, sizeof(float));
    SCUDirArg rcv(&receive_buf, dir[ 2*dim+1 ],
              SCU_REC, sizeof(float));
    // -----
    // Perform the transfers
    // -----
    SCUTrans(&send);
    SCUTrans(&rcv);
    SCUTransComplete();
    // -----
    // Keep the maximum of what you had and what
    // you just received
    // -----
    max_buf = max(max_buf, receive_buf) ;
    // -----
    // Pass on the received data
    // -----
    transmit_buf = receive_buf;
// -----
// Store the global max
// -----
*float_p = max_buf;
```

}

}

The global minimum and sum routines would follow a similar pattern.

4.6 A simple collective communications library

In a manner similar to the *hello_world* program a simple global communications library is available on the QCDSP (once Bob puts it in place). It can be found in /qcdsp/sfw/qos.5.3.3/example/glb.This directory has subdirectories include, lib, src and test. There is also a makefile in this directory of which more later.

The include directory contains the file glb.h which defines the global communication subroutines. This has to be included in any user programs using the #include directive. The src subdirectory contains the source files for the library plus a makefile to build the libraries. The lib directory is where the compiled library gets placed after it is built in the src subdirectory. This is the library that has to be linked with the user code and is called glblib.olb on the QCDSP. Finally, the test directory contains a simple program that uses the library.

To build the library and the tests, go to the toplevel directory and type make all. Once this process is finished, you should find glblib.olb in the lib subdirectory and an executable called glb_test.out in the test subdirectory.

The Library Routines

The library provides the following routines.

- void glb_sum(float * float_p): This routine performs a global sum over all the processors. On entry, float_p should point to each node's own contribution to the sum. On exit float_p will point to the result of the global sum. The algorithm for the sum is as detailed in the last section. The sum is actually performed in 64 bits of precision through a user defined type. This aspect of the operation is completely hidden from the user. At the end of the operation the result is rounded back to a 32bit result.
- void glb_max(float * float_p) : This routine finds the global maximum across
 all the processors. On entry float_p points to the each processor's own data. On exit
 float_p points to the maximum of these elements. This subroutine was listed explicitly in
 the last section.
- void glb_min(float * float_p): This routine finds the global minimum across all
 the processors. On entry float_p points to each processor's own data. On exit float_p
 points to the minimum of these elements.
- void glb_bcast(float * float_p, int root_id): This routine broadcasts the
 number pointed to by float_p on the node with unique ID root_id (as determined from

UniqueID()) to all the processors. The broadcast is implemented as a global sum with node root_id contributing the number to be broadcast at the end of float_p and all the other nodes contributing 0 to the sum.

We recommend that the reader takes a look at the source code to be found in the src subdirectory to look at the code for the above library routines. However most of them are similar to the global maximum routine listed earlier.

Using the Library

We list below the test program from the test subdirectory which shows all of the library routines in use. The program first sums the unique IDs of all the nodes, then finds the maximum and minimum of these. Finally it broadcasts the ID's of each processor in turn and sums them accross the processor grid. The code is shown below:

```
#include <stdio.h>
#include <stdlib.h>
#include <sysfunc.h>
#include <glb.h>
int main(int argc, char *argv[])
 float sum serial = 0;
 int i;
 // -----
 // Serial sum of processor ID's
 // -----
 for(i = 0; i < NumNodes(); i++) {</pre>
  sum_serial += (float) i;
 // -----
 // Put Unique ID into buffer
 // -----
 static float my_buf = (float)UniqueID();
 // -----
 // Sum Unique IDs
 // -----
 glb_sum(&my_buf);
 // -----
 // Output results
 // -----
```

```
printf("Serial Sum of IDs = %g Global Sum = %g\n", sum_serial, my_buf);
// -----
// Find Minimum ID
// -----
my_buf = (float)UniqueID();
glb_min(&my_buf);
printf("Minimum Unique ID should be 0. It is %g\n", my_buf);
// -----
// Find Maximum ID
// -----
my_buf = (float)UniqueID();
glb_max(&my_buf);
printf("Maximum Unique ID should be %g. It is %g\n",
     (float)(NumNodes()-1), my_buf);
// -----
// Do some broadcasting
// -----
for ( i = 0; i < NumNodes(); i++ ) {
 // -----
 // Process I will broadcast its ID
 // to everyone
 // -----
 if( UniqueID() == i ) {
  my_buf = (float)i;
 }
 else {
  my_buf =(float) 200;
 // -----
 // Broadcast with proc i as root:
 // Should end up with i on each node
 // -----
 glb_bcast(&my_buf,i);
 // -----
 // Do a global sum on the result
 // (should be NumElems()*i)
 // -----
```

When printing the result of each operation the program also prints out what it expects the real result to be.

Why not compile up the library as indicated before and try out the test program. On q_1, which is a 64 node development board I got the following output.

```
Serial Sum of IDs = 2016 Global Sum = 2016
Minimum Unique ID should be 0. It is 0
Maximum Unique ID should be 63. It is 63
Global sum of buffers should be 0. It is 0
Global sum of buffers should be 64. It is 64
Global sum of buffers should be 128. It is 128
.
```

Using the Library in your own programs

To use the library in your own programs you must do two things.

- Firstly, you must include the header file glb.h in every source file that uses the routines.
- Secondly, you must link to the library glblib.olb. This may involve having to edit the Makefile. If you dislike this idea (its not that bad really) simply copy the files ending in .C from the src subdirectory into the directory you are working in. The default makefiles should compile them up for you.

We shall say discuss compilation and linking in more detail in a later chapter. For now suffice it to say that both are done using the topp command.

To include the glb.h header file, it is perhaps easiest if you leave it in some sensible (Perhaps Bob will leave it in a sensible place where it can stay forever). You can then do one of two things.

• You can include it with the directive #include "sensible_path/glb.h" where for sensible_path, you should substitute the path of the sensible place where the header file lives.

- You can include it with the directive #include <glb.h>. In this case you have to make sure that the compiler will look in the sensible place where the glb.h file is by default. You can specify a list of directories for the compiler to search with the -i compiler flag.
- If you feel really unhappy about either of the above options you can always copy the glb.h file into the directory that you are working in and just include it as #include "glb.h". This way you need neither edit the makefile nor mess with compiler flags. Its the chicken way out tho' and it may leave you open to having different versions of the header file in different projects. A real nightmare...

I prefer the latter method personally, as then I only have to type the name of the directory in which the file lives once, in the makefile.

However I do repeat: If you fear and loath makefiles, do not panic for we shall clarify them in a later chapter. Until then, just copy the files from the src subdirectory and glb.h to the directory where you are coding and use the default makefiles.

4.7 Summary of Chapter

In this chapter, we have discussed collective communications. After outlining some of the most common global communication techniques, the barrier, the broadcast, gather, scatter and reduction operations we discussed which of these may be implemented on the QCDSP. We then outlined some simple algorithms for performing global reduction operations on a mesh parallel architecture. Finally, we discussed a simple library for the QCDSP which defines global sums, maxima and minima for floating point numbers.

4.8 Exercise

Can you think of more efficient ways of doing global sums on the QCDSP? You should be able to get some ideas from the QCDSP WWW page: http://www.phys.columbia.edu/cqft/qcdsp.htm where they describe the hardware support for global sums. You can also get some ideas from the book "Introduction to Parallel Computing" by Vipin and Kumar who have gone to great lengths to describe network topologies and routing algorithms in use in general parallel computers.

Chapter 5

I/O on the QCDSP

5.1 Introduction

In this chapter we discuss the facilities afforded by QCDSP to perform input and output (I/O). We consider three aspects: parallel I/O. Serial File I/O and I/O to the terminal screen during job runtime.

5.2 Parallel I/O

The QCDSP is a massively parallel supercomputer. Hence there is a high likelihood (in the case of Lattice QCD it is virtually a certainty) that data is in some way distributed amongst the numerous QCDSP nodes. The question then is, how to collect data from the nodes of the QCDSP so that it can be written to disk, or alternatively, how to get data from a disk onto the nodes of QCDSP.

A complication is that on the QCDSP only one processor, to which we shall refer hereafter as the root node (Motherboard 0, Daughterboard 0, Unique ID=0) is allowed to actually perform I/O onto a disk on the host computer. The other nodes can output to what in a UNIX system would be the *standard output* stream. This however is merely a storage buffer on the memory of the node in question, that has to be retrieved later using the **qprintf** command (as discussed in section one). In fact only node 0 can write to the output stream during run time in a way that it is echoed on the screen of a user or can be redirected into a file. The QOS has no concept of other streams such as the UNIX *standard input* and *standard error*. On the other hand the **qrun** command does allow the passing of *command line arguments* to executing programs as will be demonstrated later.

Hence, one would think that the best way to read a serial file and distribute its data amongst processors is for the root node to read the data and perform various scatter operations to distribute the data. Alternatively, saving of data could be done by the root processor gathering data from the other processors and then writing it out to disk. This is all very well, as long as the root processor can hold the whole overall dataset in memory. However, a the amount of DRAM memory on a

single node is quite small, a mere 0.5Mwords¹. Of this the run time operating system occupies 128Kwords leaving a mere 384Kwords to the user. This has to include the user code as well as the user data. Hence one would be reduced to communicating the data from the nodes to the root node, at a rate of one pair of processors at a time. Since the QCDSP has no general point to point communication routines, this process would have to be implemented using the usual nearest neighbour communications.

Another alternative would be to use a parallel file system. There is such a system currently under development for the QCDSP, however it is not yet an integral part of the QOS, and hence it shall not be discussed any further here.

The question then remains: How can parallel data be loaded and distributed in a straightforward manner on the QCDSP, and how can the data be saved from amongst all the processors into one file. On the QCDSP this is done via the **qread** and **qload** commands using objects called Node Tagged Tiles (NTFs). An NTF is a single serial file written in ASCII (can we do binary Bob?) that contains the relevant data from all the nodes of the QCDSP. Preceeding the data from a given node is a tag identifying the node which is to host the data. We will discuss their detailed. One can then develop workstation tools to convert between these NTFs and regular serial files. For those readers who intend to load and save lattices to and from the Columbia Physics environment (stored in the Gauge Connection Archive Format²) QOS provides the more specialised commands: **qload_lattice** and **qunload_lattice**.

5.3 Parallel I/O with NTFs, gread and gload

When a QCDSP program exits and the QOS returns control to the front end, the memory space of the program prior to exit is left untouched. The **qread** and **qunload_lattice** routines access this memory directly and dump a user specified number of data blocks from a user specified memory address to a file. The **qread** command can dump either to a screen or to a NTF. Its output is always in ASCII format. The **qunload_gauge** command writes directly to a file in the Gauge Connection Archiver format.

The **qload** command can be used to place data into processor memory prior to execution. It works in two ways:

• One can specify a memory address and the value to be placed there as command line arguments to **qload**. The **qload** command will then place the given value into the given memory address. The user program can then be run. It can be passed the address of the data and the number of data blocks as command line parameters. Once running the code can set a pointer to the given address and copy the relevant data to some of its own memory.

¹The wordsize of the DSP is 32bits. The physical size of a word in DRAM is 40bits of which 8 are reserved for error correction. Hence 0.5Mwords corresponds to 2Mb of useable memory. However since everything is word size for the DSP from characters to floating point numbers it makes more sense to discuss memory in terms of words than bytes.

²See the Gauge Connection Web site http://qcd.nersc.gov

• The qload command can be given the name of an NTF. The NTF contains the starting address and the number of blocks for the data, as well as the data itself, for each processor. The program can access the data by setting a pointer to the address. which can be passed as a command line argument.

The **qload_lattice** command works similarly to the **qload** command. However the gauge connection format lattice file is not node tagged. The user passes the address where the first lattice site is to be placed on each processor, and the name of the lattice file to the **qload_lattice** command. The lattice sites are then placed into processor memory. A lattice object can then be instantiated giving the address of the first lattice element to the lattice object constructor.

Of course in both the above cases dealing with **qload**, the data may be copied directly into the program data area. This can be made to be safe, as the linker produces resolved code. This means that the code produced by the linker will not be relocated afterwards at the start of execution. Hence one can print out the address of a declared array in the program, safe in the knowledge that the next time the program will run, the array will have the same address. **WARNING: This behaviour is peculiar to the QCDSP.** Furthermore one can glean information about the addresses of symbols after the compilation and linking by looking through the linker map file. We shall discuss this in more detail in the next chapter.

Before describing the details of the NTFs, and giving examples of using the **qload** and **qread** lattices we should say a few words about the organisation of memory on each node of the QCDSP. Even before that here is possibly our most severe **WARNING:** The command qload can (and will) if incorrectly used, overwrite memory used by the QOS operating system. This can have UNPREDICTABLE results ranging from no effect, through subtle bugs, to system crashes. Be careful with it!!!.

The DSP memory Map Part 1

The complete memory map of the DSP is quite complicated. In this section we shall just describe a small part of it, that pertains to compiling and running programs.

The memory addresses on a single QCDSP node can range from 000000 to FFFFFF (remember your hexadecimal numbers?). This space however is not all memory and some of it is not accessible. In particular the 0.5Mwords of DRAM are mapped between addresses 001000 to 07FFFF. The addresses from 000000 to 001000 are not accessible by the DSP in the mode it is used. Memory addresses 080000 and upwards are special. They either contain views of the data in real memory (ie that between 000000 and 07FFFF) for use by other hardware such as the circular buffer, or they contain the memory mapped images of the DSP internal memory and control registers. We may deal with some of these latter in the chapter on optimisation.

The run time operating system starts at 060000 and can stretch until 07FFFF. Consequently **only** the memory area from address 001000 to 05FFFF (inclusive) is available for user programs.

This rough sketch of the memory map is illustrated in figure 5.1.

Address Range	Function	
D00000 – FFFFFF	Used on node 0. Unused elsewhere. Generally reading	
D00000 - 111111	from or writing to here will cause the DSP to hang	
880000 – CFFFFF	Various images of the 0.5Mwords of DRAM between	
000000 - CITTIT	000000 – 07FFFF for use of the circular buffer	
820000 – 87FFFF Not used		
810000 – 81FFFF	Addresses of registers for controlling the NGA	
800000 – 80FFFF Inernal DSP Memory and control register area		
080000 – 7FFFFF	Images of 0.5Mwords of DRAM between 000000 – 07FFFF	
	for use of the circular buffer	
060000 – 07FFFF	128Kwords of DRAM reserved for Operating system	
001000 – 05FFFF	384Kwords of DRAM for User programs	
000000 – 000FFF	Inaccessible	

Figure 5.1: Memory Map of a DSP node. The Memory is word addressed.

Program Images

If you have been using the default Makefiles you should find that usually a file is produced with the .map suffix alongside your executable. In the case of the "Hello Wold" program this file would be called something like hello_world.out. (The default makefile always names output files after the directory in which the source files live).

The map file is quite useful as it tells you about how much memory your program has allocated statically during compilation, allowing you to figure out how much space you have to *qload* your NTF into.

Your program, as far as the linker is concerned consists of 6 parts of which 3 are negligible. The important three parts go by the names of *data segment*, *binary stack segment* (BSS) and the *text segment*. The unimportant parts are there mostly to supply information to debugging tools. In the case of the Tartan Linker these are called the *debug directives* segment, the *debug strings* segment and the *debug source location* segment. You need not care about these right now. Unless you switch on the option on the compiler to produce debugging information, these latter three segments will be empty.

The BSS contains amongst other things your runtime stack, on which C++ allocates automatic variables. The data segment is used to hold various constant data items and the text segment holds the actual machine level instructions of your program. Any free memory outside these segments is unallocated and should be placed onto the heap when your program starts running so that you can allocate memory from it dynamically.

The map file tells you about how your program is laid out in memory. It starts off by telling you how much space is allocated to the various segments. For example here are some excerpts from my **hello_world.map** file

Tartan Linker, SPARC/C40, v5.1.0 4/4/100 13:33:37 Copyright (c) 1986-1992 by Tartan, Inc., All Rights Reserved

Page

1

Allocation to Output Section ".bss" in module *unnamed* in file "/homeqs0/bj/hello_world/hello_world.outtof"

Offset	Length	Input Section
1000	162	"TL.Init" in module "main" in file "/homeqs0/bj/hello_world/main.tof"
1162	1	".bss" in module "qcdsp_tcroot" in file
		"/usr/local/tartan/v2.1/etc/qcdsp_v5.3.3/tcrt30bs.olb"
1163	21	"OWN" in module "inifin" in file
		"/usr/local/tartan/v2.1/etc/qcdsp_v5.3.3/tcrt30bs.olb"
1184	1000	".stack" in file
		"/usr/local/tartan/v2.1/etc/qcdsp_v5.3.3/link.lcf"

Total Allocation = 1184 (hex)

Allocation to Output Section ".data" in module *unnamed* in file "/homeqs0/bj/hello_world/hello_world.outtof"

Offset	Length	Input Section
2184	D	"DEFALT" in module "qcdsp_tcroot" in file
2191	16E	"/usr/local/tartan/v2.1/etc/qcdsp_v5.3.3/tcrt30bs.olb" ".cinit" in file
2171	101	"/usr/local/tartan/v2.1/etc/qcdsp_v5.3.3/link.lcf"

Total Allocation = 17B (hex)

Allocation to Output Section ".text" in module *unnamed* in file "/homeqs0/bj/hello_world/hello_world.outtof"

Offset	Length	Input Section
22FF	6C	" main" in module "main" in file
22F F	00	_main in module main in life "/homeqs0/bj/hello_world/main.tof"
236B	2A	".text" in module "sysfunc" in file
		"/usr/local/tartan/v2.1/etc/qcdsp_v5.3.3/tcio30bs.olb"
2395	56	".text" in module "qcdsp_tcroot" in file
		"/usr/local/tartan/v2.1/etc/qcdsp_v5.3.3/tcrt30bs.olb"
23EB	26	"dbgrpc" in module "tcrpc" in file
		"/usr/local/tartan/v2.1/etc/qcdsp_v5.3.3/tcrt30bs.olb"
2411	7	".text" in module "tcrpc" in file
		"/usr/local/tartan/v2.1/etc/qcdsp_v5.3.3/tcrt30bs.olb"
2418	32	".text" in module "tcinit" in file

You can clearly identify the three sections for the BSS, data and text segments. (They are referred to as output sections .bss, .data and .text.)

Following all this information is a summary of the output sections that tells you the total amount of memory allocated to each individual section. In my **hello_world.map** file this looks like:

Output Sections

```
Section List for Module *unnamed* in file
"/homeqs0/bj/hello_world/hello_world.outtof"
```

Number	Physical	Start	Length	Kind	Access	Name
1		1000	1184	data		".bss"
2		2184	17B	constant		".data"
3		22FF	16F	code		".text"
4		0	0	debug		"debug_directives"
5		0	0	debugstring		"debug_strings"
6		0	0	line number		"debug_source_location"
}						

This is perhaps the most useful piece of information in the map file. It tells me the starting addresses of each of my program segments and how long they are (the columns Start and Length). I can see from here that my program uses memory from address 001000 to 0022FF + 16F = 00246E inclusive. Hence if this program were to take any input via $qload\ I$ would know that it should be safe to fill up memory from 0246F upwards.

The remainder of the map file lists every single symbol used in your program. You may find it amusing to see how few of them you recognise. Here are some symbols from the **hello_world.map** file that you may recognise

```
72
                3
                       2371
                                       "_CoorZ"
          93
                3
                       2392
                                       "_CurrentStatus___FiPc"
          97
                3
W
                       2396
                                       " c int00"
          6D
                3
                       236C
                                       " DbNum"
         14C
                3
                       244B
                                       " exit"
         14B
                3
                       244A
                                       "_exit$LAJ"
          8E
                3
                       238D
                                       "_fclose"
          8C
                3
                                       "_fopen"
                       238B
                3
          8D
                                       "_fprintf"
                       238C
          91
                3
                       2390
                                       "_HdwCheck___Fi"
          94
                3
                                       " InterruptExit FiPc"
                       2393
           1
                3
                       2300
                                       " main"
           0
                3
                                       "_main$LAJ"
                       22FF
                3
          6C
                       236B
                                       " MbNum"
          90
                3
                       238F
                                       " NodeStatUpdate Fv"
                                       "_NumNodes"
          77
                3
                       2376
          7C
                3
                                       "_printf"
                       237B
```

You may be worried here that at startup, spurious memory allocation might overwrite your qloaded data. This is not a worry here as the map file includes the heap area as well. Hence if you place something above the last address in the map file, it is guaranteed that it will not be malloced over.

You can configure the size of the heap and other program layout options in a file called the Linker Control File. If you do not use one of these the program will be linked under the control of a default linker control file which allocates a pitifully small heap. (0x1000 words). Hence when you come to code bigger projects you will have to allocate the heap yourself in your own LCF. This will involve a partitioning of the available memory between room for your program, data you intend to allocate dynamically (heap) and data you intend to gload.

The standard procedure is to qload data very high in memory, just below the operating system. Suppose you have data of say 32Kwords per processor (0x8000 words (hex)). The operating system starts at address 0x060000 and so you might want to load the data to 0x06000 - 0x008000 = 0x058000. Then you can set up your LCF so that your program cannot extend beyond 0x057fff (heap included). It is of course prudent to leave a bit of room here and there in case you modify your program. In the words of investment brokers "The size of your program can go down as well as up.".

We will discuss linking, program sections and control files in more detail in the next chapter about compiling and linking.

Node Tagged File Formats

We now discuss node tagged files. Node tagged files are written in ASCII and contain hexadecimal numbers. Although this can cause the files to be large and slow to read and write, it does have the advantage that the files are readable without needing specialist dump programs.

The file format consists of one hexadecimal entry on each line written in ASCII. The structure of the entries is given below. Each heading and sub-heading corresponds to an entry in the NFT.

- **A 'Magic Number':** This is an integer for identifying which of the 3 types of node tagged files one is dealing with. The three types are:
 - tree the nodes are tagged according to their position on the SCSI tree by a pair of numbers (M,D) representing the node's motherboard and daughterboard number respectively. Each of these numbers corresponds to a new entry on a new line in the NTF. The magic number for this mode is 2. This is the 'default' tag mode. Node tagged files in this format usually have a **.tree** file suffix.
 - Machine 4D the nodes are tagged with their 4D coordinates within the processor grid (A four-tuple of integers, each of these on a new line in the NTF). The magic number for this tag mode is 3. Node tagged files in this mode usually have a **.m4d** file suffix.
 - Physics 4D to all intents and purposes this format the same as the machine 4D. Nodes are tagged by their 4D coordinates. The meaning of fields however is different. In terms of machine coordinates, everything is fixed. However the machine coordinates can be re mapped at boot time to interchange the order of coordinates. Hence (T,X,Y,Z) coordinates can be re mapped to (X,Y,Z,T) coordinates etc. The machine 4D file never changes the data layout. However the physics 4D file can shift the data around depending on the current 4D mapping. The magic number for this mode is 4. Node tagged files in this mode usually have the **.p4d** file suffix.
- **A sequence of Node Records** for each node in the processor there is a record. A record has the following structure.
 - A Node Tag A tag identifying the node. For the tree structure this consists of the motherboard number followed by the daughterboard number. For the 4D modes this is a four tuple of integers identifying the 4D coordinates of the node.
 - The number of blocks for the Node An integer indicating how many blocks of data are to be loaded into the memory of the node (Hexadecimal without leading 0x)
 - Base Address A hexadecimal address (without leading 0x) specifying the location in memory where the first data item is to be placed.
 - Data Blocks An ASCII representation of each data block. In the case of integers this is quite readable. In the case of floats it is not
 - A terminator An ASCII zero follows the final record as a record termination signal.

The formats of the NTFs, the magic numbers, node tags and file extensions are summarised in figures 5.2, 5.3 and 5.4.

Perhaps the easiest way of creating a node tagged file is to dump some memory from a program. We now proceed to demonstrate how this is done.

Magic Number
Record No 1
Record No 2
:
Last Record

Figure 5.2: Structure of a Node Tagged file – Components are sequential in an ASCII file. Each component starts on a new line

Node Tag
n – Number of blocks
Base Address
block[0]
block[1]
:
block[<i>n</i> − 1]
0 (terminator)

Figure 5.3: Structure of a Node tag record – Each component is in ASCII on a new line

5.4 Creating Node Tagged files using qread

Consider the following little program:

The code itself does nothing, but prints out the address of the first element of the floating point number array i. We can use this information to dump the floating point array into an NTF.

File Type	Magic Number	umber Node Tag Components	
Tree	2	Motherboard No	.tree
1166	2	Daughterboard No	
		T Coordinate	
Machine 4D	3	X Coordinate	4 -7
		Y Coordinate	.m4d
		Z Coordinate	
		T Coordinate	
Physics 4D	4	X Coordinate	
		Y Coordinate	.p4d
		Z Coordinate	

Figure 5.4: Magic Number Node Tag and File Suffix information for the three node tagged file types. The magic number is the first number of any node tagged file. It should be an ASCII integer on the first line. The node tag fields should be ASCII integers. Each entry should take a separate line in the file.

Setting options for gread

There are several options that allow you to control how you want to perform the dump. For example, you can dump to a screen or a named file. You can dump into an NTF of either the tree, machine 4D or of the physics 4D variety. You can rewrite an existing file or append to it. These options are set by the following qcsh commands:

qset_read_output_filename: This command takes one argument which is the name of the file you want qread to dump your data to.

qset_read_output_file_access: This command takes a single argument to determine whether you wish to use the append write mode or the overwrite mode. The possible arguments are

a – append mode

w - write mode

qset_read_output_select : This command determines whether you wish your dump to
 go to the screen or to a named file (that you have to name using qset_read_output_filename.
 It takes a single argument. The argument can have the following values:

f – dump to named file

s – dump to screen

qset_read_output_tagged: This command determines the tag type of your dump file. It takes a single argument. Allowed argument values are

no – Do not node tag the file. I do not describe the results of this option.

```
    yes - Place node tags into file. Use default tag mode (tree).
    tree - Use tree tag mode.
    m4d - Use machine 4D tag mode.
    p4d - Use physics 4D tag mode.
```

Using qread to dump to a file

Let us now have a go at dumping the floating point array in the previous code snippet. Create a directory called dump somewhere and copy the program listing above into a file named main.C (alternatively the directory may be available under <QOS_VERS>/examples/ParIO/dump).

If you are doing things by hand, you probably want to copy a default Makefile from the hello_world directory and build the program by typing make. This should produce a program called dump.out. If you are copying the directory structure it should have a Makefile with it.

Boot the machine and enter the QC-Shell.I will use the machine q_1. Now let us suppose that we want to dump the float array from each PE into a file called foo. Let us try first with the tree node tag mode. We set up the dump using the following qcsh commands (I removed the Qdaemon responses to save space).

```
(dump: qcsh[q_1])% qset_read_output_select f
(dump: qcsh[q_1])% qset_read_output_filename foo.tree
(dump: qcsh[q_1])% qset_read_output_file_access w
(dump: qcsh[q_1])% qset_read_output_file_tagged tree
```

Now run the program using grun as before (Qdaemon messages removed):

```
(dump: qcsh[q_1])% qrun dump.out
The address of i is 105d and its length is 5 blocks
```

Now I can dump to the NTF using <code>qread</code>. This command takes the base memory address and the number of blocks as two hexadecimal numbers respectively (without the leading 0x in terms of notation). I dump the array with the command (of course you may find a different base address to 105d in which case you have to substitute that in what follows):

```
(dump: qcsh[q_1])% qread 105d 5
```

to which the machine responds:

```
Qdaemon state is:
qdaemon task is READ
```

```
qdaemon abort no and resume no
       QCDSP synchronous
       All nodes selected with SCSI tree coordinates
Qdaemon: Parameters for read from QCDSP:
       Read address:
                               0x105d
       Read blocksize:
                               0x5
       Filename:
                              foo.tree
       File access:
                              write
       File format:
                               node tagged tree
Odaemon state is:
       qdaemon task is READ
       qdaemon abort no and resume no
       QCDSP synchronous
       All nodes selected with SCSI tree coordinates
```

and the file foo.tree magically appears in the directory.

We can actually look at the **foo.tree** file. My one looks like:

```
2
0
0
5
105d
1000000
1400000
2000000
2200000
0
0
1
105d
0
1000000
1400000
2000000
2200000
```

```
0
3f
5
105d
0
1000000
1400000
2000000
0
```

You can clearly see the structure here. The first number in the file is the magic number 2. The next two are the node tag for the tree coordinates. These are motherboard and daughterboard identifiers (the last daughterboard identifier is 3f which is hex for 63. This is because I am using only a 64 PE single motherboard machine). After this you have the number of blocks -5 – followed by the base address 105d. You have 5 lines of data. They are the hex representations of the bit pattern that makes up our 5 floating point numbers. Finally you have a 0 as an end of record marker.

5.5 Loading data with qload

Poking individual words with qload

The command <code>qload</code> allows you to either load node tagged files or to place individual words into memory. Like qread it also has a number of options you can set. In this subsection we discuss first how to place individual words in memory (ever hear of the POKE keyword in BASIC? This is kind of the same stuff).

The incantation to load a word into a particular location in memory is of the format

```
gload <address> <data>
```

where both the address and the data are hexadecimal values. For example to load the value 5 into the memories of all the processors at address 04EFFF I would type a magic incantation like:

```
(qcdhost/homeqs0/bj/QCDSP/IO/load: qcsh[q_1])% qload 04EFFF 5
```

To that this has worked consider the following program:

```
#include <stdio.h>
#include <stdlib.h>
#include <strings.h>
#include <sysfunc.h>
int main( int argc, char *argv[] )
{
```

```
// qload will put (poke?) the distributed memory
// into the data. User passes in the address the
// data has been load to through command line
// arguments
int *loaded_ints;
int nblocks;
// Check command line arguments
// -----
if( argc != 3) {
 printf("Usage: qrun load.out <base address(in hex)> <no of blocks(in hex)>\n
 return(EXIT_FAILURE);
// Convert HEX base address string into an integer pointer value
// and point loaded_ints there
// -----
loaded_ints = (int *)strtol(argv[1], (char **)NULL, 16);
// Convert no_of_blocks string into an integer value
nblocks = (int)strtol(argv[2], (char **)NULL, 16);
// -----
// Confirm values to the user
printf("%x (%d) blocks of data at %x\n", nblocks, nblocks,
     (unsigned int)loaded ints);
// Print back the data
// -----
for(int i = 0; i < nblocks; i++) {
 printf("%d ",loaded_ints[i]);
printf("\n");
return(EXIT_SUCCESS);
```

This program basically checks your qload. You need to pass it the base address to where you have loaded your data as well as the number of blocks comprising the data. These are passed on to the program as strings in the argv array, so first we have to convert them to addresses. This is done using the strtol function call which converts a string to a long integer.

The number 16 in the third argument tells strtol³ that the string is supposed to represent a hexadecimal number. We set a pointer loaded_ints to the results of the string conversion, in other words to our base address.

We then also need to convert the number of blocks to an integer. This is done by the second call to strtol.

After confirming to the user that the input values have been read correctly, the code proceeds to

³strtol is a very useful function. For details look at the man page in section 3 of the manual

list the specified number of words from the base address, treating them as integers and prints them to the screen.

Compile up this little program. You should also find it already written on the QCDSP in the directory <QOS_VERS>/examples/ParIO/load with a suitable default Makefile (Replace this with wherever we put it...) Now lets have a look in the map file (in my case its called load.map):

Output Sections

Section List for Module *unnamed* in file "/homegs0/bj/OCDSP/intro/IO/load/load.outtof"

Number	Physical	Start	Length	Kind	Access	Name
1		1000	110C	data		".bss"
2		210C	102	constant		".data"
3		220E	26F	code		".text"
4		0	0	debug		"debug_directives"
5		0	0	debugstring		"debug_strings"
6		0	0	line number		"debug_source_location"

This tells me that it is safe for me to load things above the address 220E+26F = 247D. However it is good practice to load things high in the memory so I will load the data to 05FFE0. I do this by typing

```
(qcdhost/homeqs0/bj/QCDSP/intro/IO/load: qcsh[q_1])% qload 05FFE0 5
```

Now I run my program

```
(qcdhost/homeqs0/bj/QCDSP/intro/IO/load: qcsh[q_1])% qrun load.out 05FFE0 1
```

The computer replies

```
1 (1) blocks of data at 5ffe0 5
```

on all the processors.

Loading NTFs using qload

The gload command can also be used to load NTFs. Since all the information about the data is in the NFT (processor coordinates, base addresses, number of blocks and the actual data) is present in the NTF file. There is no need to specify these. The command is simply

```
qload <filename>
```

The command can identify the tag type used in the file from the magic number.

Consider the little program above but now converted to dump floats instead of ints:

```
#include <stdio.h>
#include <stdlib.h>
#include <strings.h>
#include <sysfunc.h>
int main( int argc, char *argv[] )
 // gload will put (poke?) the distributed memory
 // into the data. User passes in the address the
 // data has been load to through command line
 // arguments
 // -----
 float *loaded_floats;
 int nblocks;
 // Check command line arguments
 // -----
 if( argc != 3) {
   printf("Usage: qrun load.out <base address(in hex)> <no of blocks(in hex)>\n
   return(EXIT_FAILURE);
 }
 // Convert HEX base address string into an integer pointer value
 // and point loaded_ints there
 // -----
 loaded_float = (float *)strtol(argv[1], (char **)NULL, 16);
 // Convert no_of_blocks string into an integer value
 nblocks = (int)strtol(argv[2], (char **)NULL, 16);
 // Confirm values to the user
 printf("%x (%d) blocks of data at %x\n", nblocks, nblocks,
       (unsigned int)loaded_ints);
 // Print back the data
 // -----
 for(int i = 0; i < nblocks; i++) {</pre>
   printf("%f ",loaded_floats[i]);
 printf("\n");
 return(EXIT_SUCCESS);
```

We can use this to load our floating point NFT (foo.tree) that we created before, but first another word about memory.

Manual Data Relocation

Usually it will be the case, that a program loads an NTF at the start to a memory location, then processes it and then dumps it at the end. In this case the address in the node tag file is fine both for reading and dumping. However we are now at an advanced stage, where one program produced

the NTF but another program wants to load it.

In the program that made the NFT, the data was in a standard statically defined array that started (in my case) at address 105d. Now however, we want to load it into a different program. After converting the dump program to deal with floats, a quick look at the map file tells me that memory is still free from above 247D – this is a consequence of the fact that the int on the DSP is the same size as the float. It is desirable to load the data to a high address, say 05e000. This can be achieved by editing the NTF and replacing 105d everywhere with 05e000. In my case the it is simply a question of a global edit. This is one of the advantages of the NTF being in ASCII format.

Example of loading an NTF

So, edit the previous program to deal with floats as I have done above. Edit your file tree.foo replacing the base addresses everywhere with 05e0000. Load the NTF using the command

```
\label{load:qcdhost/homeqs0/bj/QCDSP/intro/IO/load: qcsh[q_1]) % qload foo.tree $$ and run the program, it should probably be called load.out: $$ $$
```

```
(\verb|qcdhost/homeqs0/bj/QCDSP/intro/IO/load: \verb|qcsh[q_1]|) % | qrun load.out | 05e000 | 5e000 | 5e000 | for the content of the
```

The computer should reply:

```
5 (5) blocks of data at 5e000
1.000000 2.000000 3.000000 4.000000 5.000000
```

on all the processors.

Setting up qload options

You should now be able to use <code>qload</code> to load parallel data. However just like the case of <code>qread</code> there are some options that you can set. These allow you to specify the filename to <code>qload</code> or simply just the address and value to <code>qload</code> (for poking single data). Then you can use <code>qload</code> with no arguments at all. The commands to set the options are

qset_load_select: Specifies whether qload should read from a file or from the screen. The command can have up to 1 argument whose value is either file or screen (alternatively f or s respectively). If no arguments are given, the current value of the option is shown. I generally found that this option reset itself to file every time qload is invoked with an explicit filename.

qset_load_screen: Specifies an address and a data item. This option sets things up so that a subsequent call to qload will load the specified data to the specified address. The command can have either two arguments or none. Invoking the command with no arguments causes it to print the current values for the address and the data value. When the command is invoked with two arguments, the meaning of the arguments are

```
address specifies the address to load to.

data value specifies the value of the data to be loaded.
```

Notes:

- The following are equivalent: qload <address> <value> and qset_load_select s; qset_load_screen <address> <value>; qload.
- If one selects to load from the screen, and sets the address and value and then invokes qload <filename>, then qload will switch to reading from the file instead. If one has set a filename to load from using qset_load_filename and has switched qload to read from file using qset_load_filename file, if one now executes qload <address> <value> then qload will switch to screen mode and the last option will be valid.

qset_load_filename: This command sets the name of the file that qload will use to load data from. It can have either one or no arguments. If the command is invoked with no arguments it will echo the current value of the filename. If invoked with one argument the argument, the argument should be the name of the file to load from. Note: The following are equivalent: qload <file> and qset_load_select file; qset_load_filename <file>; qload.

5.6 Digression for Lattice Folk

There are two specialised commands to load and unload lattice gauge fields to and from the QCDSP. These are, as mentioned previously qload_lattice and qunload_lattice. Before describing these I should say a few words about the lattice file format.

Gauge Connection Format

Currently NERSC provides an archive area called the Gauge Connection. They use a particular format for lattices there. This format consists of a header section and a data section. The header section consists of a whole lot of fields. This is useful for keeping track of where the lattices came from. The header file also keeps information about the average plaquette to assist in validation.

The command <code>qload_lattice</code> simply takes a configuration in the NERSC archive and loads it into memory at a user supplied starting address. For example, prior to running a certain piece of code. I have obtained (actually output from another program) a lattice called <code>lattice.ieee.full</code>

. The .ieee part of the naming convention indicates that the floating point format of the lattice is IEEE 32Bit big—endian. The .full part indicates that all the link matrices are stored in full (as opposed to compressed 2 row format).

I would like to load this lattice prior to computation. It has a volume of 8^4 sites, which is to be distributed over 64 processors. My processor grid is a $4 \times 2 \times 2 \times 4$ lattice (the order of the dimensions in this case is $X \times Y \times Z \times T$. This implies that the local lattice size per processor is $2 \times 4 \times 4 \times 2 = 64$ sites. On each site there are 4 link matrices each of these being a full SU(3) matrix consisting of 9 complex numbers which are represented as 18 real numbers. Hence the size of the local data is $64 \times 4 \times 18 = 4608$ numbers and since each number is word sized this means that the data requires 4608 words of storage which in hexadecimal is 1200. Thus the highest I can load this data in memory is to address 060000 - 001200 = 5ee00. To be a little conservative I choose to load it to 5ed00.

The instruction I give to the QCDSP is:

```
(qcdhost/homeqs0/bj/QCDSP/sfw/hello_world: qcsh[q_1])% qload_lattice 5ed00
```

Apart from the usual qdaemon messges, the response of the computer was as follows:

```
DoLoadLattice: a copy of the header from lattice.ieee.full
       BEGIN HEADER
       CHECKSUM = 1649a43
       LINK_TRACE = +0.0020214
       PLAQUETTE = +0.5132502
       DATATYPE = 4D_SU3_GAUGE_3x3
       HDR VERSION = 1.0
       STORAGE_FORMAT = 1.0
       DIMENSION_1 = 8
       DIMENSION_2 = 8
       DIMENSION_3 = 8
       DIMENSION 4 = 8
       BOUNDARY_1 = PERIODIC
       BOUNDARY_2 = PERIODIC
       BOUNDARY_3 = PERIODIC
       BOUNDARY_4 = ANTIPERIODIC
       ENSEMBLE_ID = BALINT_LAT_1
        ENSEMBLE_LABEL = Balint's first test lattice 8x8x8x8 pure gauge beta=6.0
       SECTIONCE NUMBER = 100
       CREATOR = Columbia
       CREATOR_HARDWARE = QCDSP
       CREATION_DATE = Fri Apr 28 15:50:26 2000
       ARCHIVE_DATE = Fri Apr 28 15:50:26 2000
       FLOATING_POINT = IEEE32BIG
       END_HEADER
DoLoadLattice: keywords used to load lattice
       CHECKSUM = 1649a43
       LINK_TRACE = 0.0020214
       PLAQUETTE = 0.5132502
       DATATYPE = 4D SU3 GAUGE 3x3
       DIMENSION_1 = 8
       DIMENSION_2 = 8
       DIMENSION 3 = 8
       DIMENSION_4 = 8
       FLOATING_POINT = IEEE32BIG
DoLoadLattice:
       Loading 4608 words per node to address 0x5ed00
```

```
Total transfer of 294912 words to 64 nodes of QCDSP
        Starting to read 294912 words from host disk
       Finished loading 1/4 of lattice to QCDSP
       Finished loading 2/4 of lattice to QCDSP
       Finished loading 3/4 of lattice to OCDSP
       Finished loading 4/4 of lattice to QCDSP
       Converting from IEEE32BIG to TIDSP32 on QCDSP
DoLoadLattice:
       Loading plag.gin to calculate plag and trace on OCDSP
Output from plaq.qin running on QCDSP
       Machine size (X,Y,Z,T) (4,2,2,4)
       Lattice size per node (2,4,4,2)
       Total lattice size (8,8,8,8)
       Lattice address 0x5ed00
                       0.5132502
       plag
        link trace
                       0.0020214
MkerExecAck: starting to check program exit status
AnalyzeQCDSP::SCUCheck: SCUDebug mode 0 doesn't allow analysis
MkerExecAck: SCUCheck could not be run
DoLoadLattice:
       Plaquette from header and QCDSP (0.5132502) agree to 0.000001
       Link trace from header and QCDSP (0.0020214) agree to 0.000001
       Checksum from header and QCDSP agree (0x1649a43)
```

The messages are quite self explanatory. First gload lattice echoed back the contents of the header file. Then it loaded the lattice, partitioned it amongst the processors. Thereafter it converted from IEEE 32Bit big endian format to the DSPs internal format denoted TIDSP32. A separate program was then run to calculate the plaquette and the average link trace to validate the data. Finally QCDSP responded indicating that the lattice has been loaded, that the plaquette and link trace have been found to be correct to 4dp and that the checksum test has been passed.

At this point one can start one's computation. To be able to use the lattice in the Columbia Physics System, one has to set the members of a global instance of a GlobalJobParameter class (it is required to have one of these called GJP declared globally). One has to set the private start_conf_kind member of the GlobalJobParameter class to have value START_CONF_LOAD, and one has to set the private start_conf_load_addr pointer to point in memory to first word of the loaded lattice.

Currently this can only be done through the public Initialize member function of the GlobalJobParamet class. This function takes as an argument a reference to a DoArg class. The DoArg class has public member variables start_conf_kind and start_conf_load_addr. The former of these has to be set to START_CONF_LOAD and the latter to the address of the first word of the loaded data

The address can be passed to the program as a command line argument, making it available to each processor, or it could be read by node 0 using the usual C style file I/O and then broadcast to the rest of the system. An example snippet of code to set up the lattice may look something like as follows:

```
// -----
// Standard headers
// -----
```

```
#include <stdio.h>
                // Exit
#include <stdlib.h>
// QCDSP Non Physics Includes
// -----
#include <sysfunc.h>
// -----
// QCDSP Physics Includes
#include <util/include/lattice.h> // Lattice Classes
#include <util/include/gjp.h>
                           // Global Job Parameters Class
#include <util/include/verbose.h> // Verbose Output
// -----
// These are always defined Globally
// -----
GlobalJobParameter GJP;
Verbose VRB;
// -----
// PE 0 is boss
// -----
#define BOSS_ID 0
main(int argc, char *argv[])
 // -----
 // Check Arguments
 // Can call with either no arguments (ordered start) => argc = 1
 // Or with one argument, being hex address of gauge field
 // -----
 if ( argc > 2 ) {
   printf("Usage: qrun prog.out [hex address (no leading 0x)]\n");
   exit(EXIT_FAILURE);
 // -----
 // Initialize the Global Job Params
 DoArg do_arg;
 int my_id = UniqueID();
 // -----
 // Global lattice {\tt X} {\tt Y} {\tt Z} {\tt T}
 // -----
 int g_latt_size[4] = { 8, 8, 8, 8 };
 // Local lattice X Y Z T
 int l_latt_size[4] = { 2, 4, 4, 2 };
 printf("Local Lattice Size is: %d %d %d %d\n",
1_latt_size[0],l_latt_size[1],l_latt_size[2], l_latt_size[3]);
 // -----
 // Use Whole Machine
```

```
// -----
do_arg.x_nodes = SizeX();
do_arg.y_nodes = SizeY();
do_arg.z_nodes = SizeZ();
do_arg.t_nodes = SizeT();
// -----
// We are in 4D by I set the 5th dimension anyway
// -----
do_arg.s_nodes = 1;
// -----
// Set the lattice volume -- Sites per node
// -----
do_arg.x_node_sites = l_latt_size[0];
do_arg.y_node_sites = l_latt_size[1];
do_arg.z_node_sites = l_latt_size[2];
do_arg.t_node_sites = l_latt_size[3];
do_arg.s_node_sites = 1;
// -----
// Set Boundary Conditions
// -----
do_arg.x_bc = BND_CND_PRD; // Periodic
do_arg.y_bc = BND_CND_PRD; // Periodic
do_arg.z_bc = BND_CND_PRD; // Periodic
do_arg.t_bc = BND_CND_APRD; // Antiperiodic
// Set lattice start depending on no of arguments
// -----
if( argc == 1 ) {
                       // No user supplied arguments
  do_arg.start_conf_kind = START_CONF_ORD;
else { // User supplied load address
  do_arg.start_conf_kind = START_CONF_LOAD;
  // Hex String to address convertion
  // -----
  do_arg.start_conf_load_addr = (Matrix *)strtol(argv[1],(char **)NULL, 16);
do_arg.start_seed_kind = START_SEED_FIXED;
do_arg.verbose_level = VERBOSE_RESULT_LEVEL + 1;
// Quenched QCD -- 3 colours, beta = 6
// -----
do_arg.colors = 3;
do_arg.beta = 6.0;
// -----
// Initialise GJP Structure
// -----
GJP.Initialize(do_arg);
// -----
// Set verbosity
VRB.Level(VERBOSE_RESULT_LEVEL+1);
// Grab a Wilson Gauge Lattice. with possibly wilson fermions in future
// This gets initialised from GJP
```

```
// -----
GwilsonFwilson lat;
// -----
// Calculate The Global Trace of the plaquette
Float t;
Float normal; // Plaquette normalization factor
normal = g_latt_size[0]*g_latt_size[1]*g_latt_size[2]*g_latt_size[3];
normal *= 6 * 3; // No of planes * 3 to normalize unit gauge to 1
t=lat.SumReTrPlaq()/normal;
// Print Plaquette
if( my_id == BOSS_ID ) {
 printf("Boss: Startup Gauge Field -- Sum Tr Plaq = %f\n",(float)t);
// -----
// Do your world beating calculation below
// -----
// Your world beating calculation ends
// You want to store gauge field perhaps
// -----
// -----
// Print Gauge Field -- for saving
printf("Local Gauge field starts at %x\n",(int)lat.GaugeField());
exit(EXIT SUCCESS);
```

Running this program with no arguments produced the following output

```
ocal Lattice Size is: 2 4 4 2

Boss: Startup Gauge Field -- Sum Tr Plaq = 1.000000

Local Gauge field starts at 137c8
```

which is the correct output a unit gauge. I could recover the gauge field by using qunload_lattice as I have printed out the address of the start of the gauge field data.

Now I load the gauge field as before (I have decided to omit the output from qload_lattice as I have already included it in full earlier. I then run the program giving the starting address (5ED00) as my first argument. I get the following reply from the QCDSP:

```
Local Lattice Size is: 2 4 4 2

Boss: Startup Gauge Field -- Sum Tr Plaq = 0.513250

Local Gauge field starts at 5ed00
```

Indicating that I managed to initialize the lattice correctly.

Unloading Gauges

Now we can try to unload the gauge and keep it. A few things need to be borne in mind. It is a local convention, that gauge fields in use here at Columbia be kept in the DSP native representation (TIDSP32) whereas fields that are to be shipped off to the gauge connection should really be in the IEEE big—endian format (IEEE32BIG).

Also, gauge fields can be kept in compressed format (2 of the 3 rows stored only) for each SU(3) matrix or in full format.

While the <code>qload_lattice</code> command can learn all this from the header part of the gauge field configuration file, and the user program has knows about the lattice through the <code>GlobalJobParameter</code> structure, the <code>qunload_lattice</code> command, running on the front end, knows nothing about our lattice and has to be told via a parameter file.

The parameter file essentially provides the header information part of the gauge configuration data file. It consists of an ordered list of entries, which we shall describe below. Each entry lives on a single line of the file. Lines containing no text are ignored. Lines starting with a # character are treated as comments and ignored. The # character does not need to be in the first column of the line. However it may not follow an entry. There are therefore no trailing comments. There should be no trailing spaces following entries. **The order of entries is important.**

The entries of the file, in order are

T Size – this is a single integer giving the global lattice size in the T (Euclidean Time direction)

X Size – this is a single integer giving the global lattice size in the X direction

Y Size – this is a single integer giving the global lattice size in the Y direction

Z Size – this is a single integer giving the global lattice size in the Z direction

Boundary condition for T – This entry specifies the boundary condition for the T direction. The boundary conditions can be either periodic or antiperiodic. This entry can take the values PERIODIC or ANTIPERIODIC respectively.

Boundary condition for X – This entry specifies the boundary condition in the X direction. The entry can take the values PERIODIC or ANTIPERIODIC.

Boundary contidion for Y – This entry specifies the boundary condition in the Y direction. The entry can take the values PERIODIC or ANTIPERIODIC.

Boundary contidion for Z – This entry specifies the boundary condition in the Z direction. The entry can take the values PERIODIC or ANTIPERIODIC

Lattice Base Address – The address in memory where the data for the lattice starts on each processor. Can be found using the GaugeField() member function of the lattice base class in the Physics system. The number should be entered in hexadecimal without any preceding 0x.

- **Lattice Output Filename** The filename where the lattice configuration is to be dumped. Be aware of the file name suffix convention.
 - .dsp.full Gauge field is stored in DSP internal floating point representation. All three rows of the link matrices are stored in the file.
 - .dsp Gauge field is stored in DSP internal floating point representation. The link matrices are stored in compressed 2 row format. The third row is removed.
 - .ieee.full Gauge field is stored in IEEE 32Bit Big Endian representation. All three rows of the link matrices are stored.
 - .ieee Gauge fields are stored in IEEE 32Bit Big Endian format. The link matrices are stored in compressed 2 row format. The third row is removed.
- **Ensemble ID** This is some label uniquely identifying the simulation. It can be any string. On the Gauge Connection, this label is displayed when your ensemble is listed. By convention configurations from Columbia have ensemble ID's of the form CU_XXXX where CU identifies Columbia University (MILC have MILC and Ohio State University use OSU) and XXXX is a unique number identifying the ensemble from which the gauge is from.
- **Ensemble Label** This should be a more human readable label giving some details about the ensemble such as the parameters that were used to generate it. This entry takes a string value.
- **Sequence Number** This is an integer value identifying the configuration in a given ensemble (for example an HMC trajectory number, or a Heat bath sweep number).
- **Creator String** A string describing the group responsible for the generation of the gauge configuration. Columbia uses Columbia.
- **Creator Hardware** The machine used to create the lattices. If the configurations were generated by QCDSP, this field should either be CU-QCDSP for Columbia or RBRC-QCDSP for the RIKEN–Brookhaven machine.
- **Datatype** An entry to describe how to save the configuration. This can take one of the following values
 - 4D_SU3_GAUGE Gauge field is to be saved in compressed 2 row format.
 - 4D_SU3_GAUGE_3x3 Gauge field is to be saved in full 3 row format.
- **Floating point type** Specifies whether the gauge should be saved in DSP native format or in IEEE format. This can take values:
 - TIDSP32 Gauge configuration to be saved in Texas Instruments DSP 32bit format.
 - IEEE32BIG Gauge configuration is to be saved in 32 bit IEEE Big Endian format.
- **Creation Date** A character string specifying the date of creation. It entry has the value now it will be timestamped automatically when qunload_gauge is invoked.

For example to dump the gauge we have just loaded in the previous section I would need the following file:

```
# size in T
# size in X
# size in Y
# size in Z
# boundary condition for T
ANTIPERIODIC
# boundary conditionf for X
PERIODIC
# boundary condition for Y
PERIODIC
# boundary condition for Z
PERIODIC
# lattice base address
05ed00
# lattice output file
lattice.ieee.full
# Ensemble ID
BALINT_LAT_1
# Ensemble label
Balint's first test lattice 8x8x8x8 pure gauge beta=6.0
# Sequence number
100
# Creator string
Columbia
```

```
# Creator Hardware
QCDSP

# Datatype
4D_SU3_GAUGE_3x3

# Floating point type
IEEE32BIG

# Creation Date
now
```

Before you can use qunload_lattice to unload a configuration you must tell it about the lattice. If you have created a file such as above you can tell qunload_lattice to use it as a description file with the QC-Shell command:

```
qset_unload_lattice f <file>
```

where <file> is the name of your description file.

Alternatively you can enter all the fields manually from the command line by typing:

```
qset_unload_lattice m
```

and enter answers to the various prompts.

There is a default mode for qunload_lattice which you can set with the command:

```
qset_unload_lattice d
```

The default settings however refer to a lattice of 2^4 sites with periodic boundaries in each direction. Hence the default mode is probably not very useful.

Once you have executed one of the qset_unload_lattice commands and have got no error messages, you are ready to unload the lattice. It is a simple matter of typing

```
qunload\_lattice
```

For example, I now try to unload the unit gauge lattice I can produce by running my gauge reading program with no arguments. I use the example lattice description I have listed above except I've changed the output filename to unit_gauge_8888.ieee.full. The description is in a file I call qunload.in.

I first run the program without arguments to produce the unit gauge

```
(qcdhost/homeqs0/bj/QCDSP/sfw/hello_world: qcsh[q_1])% qrun simple_test.out
```

I get the response

```
Local Lattice Size is: 2 4 4 2

Boss: Startup Gauge Field -- Sum Tr Plaq = 1.000000

Local Gauge field starts at 137c8
```

I edit my description file qunload.in to change the base address to 137C8. I execute the commands:

```
(qcdhost/homeqs0/bj/QCDSP/sfw/hello_world: qcsh[q_1])% qset_unload_lattice f qunload.in
(qcdhost/homeqs0/bj/QCDSP/sfw/hello_world: qcsh[q_1])% qunload_lattice
```

To which the QCDSP replies:

```
DoUnloadLattice:
        Unloading 4608 words per node from address 0x137c8
        Total transfer of 294912 words from 64 nodes of QCDSP
       Starting to write out 294912 words to host disk
       Finished writing out 1/4 of lattice
       Finished writing out 2/4 of lattice
       Finished writing out 3/4 of lattice
        Finished writing out 4/4 of lattice
       Loading plaq.qin to calculate plaq and trace on {\tt QCDSP}
Output from plaq.qin running on QCDSP
       Machine size (X,Y,Z,T) (4,2,2,4)
       Lattice size per node (2,4,4,2)
       Total lattice size (8,8,8,8)
       Lattice address 0x137c8
        plag
                        1.0000000
        link trace
                        1.0000000
MkerExecAck: starting to check program exit status
AnalyzeQCDSP::SCUCheck: SCUDebug mode 0 doesn't allow analysis
MkerExecAck: SCUCheck could not be run
DoUnloadLattice:
       copy of header written to lattice.ieee.full
        BEGIN_HEADER
       CHECKSUM =
       LINK_TRACE = +1.0000000
       PLAQUETTE = +1.000000
       DATATYPE = 4D_SU3_GAUGE_3x3
       HDR\_VERSION = 1.0
       STORAGE_FORMAT = 1.0
       DIMENSION_1 = 8
       DIMENSION_2 = 8
       DIMENSION_3 = 8
       DIMENSION_4 = 8
       BOUNDARY_1 = PERIODIC
       BOUNDARY_2 = PERIODIC
       BOUNDARY_3 = PERIODIC
       BOUNDARY_4 = ANTIPERIODIC
        ENSEMBLE_ID = BALINT_LAT_1
        ENSEMBLE_LABEL = Balint's first test lattice 8x8x8x8 pure gauge beta=6.0
        SEQUENCE_NUMBER = 100
```

CREATOR = Columbia CREATOR_HARDWARE = QCDSP CREATION_DATE = Fri Apr 28 18:29:07 2000 ARCHIVE_DATE = Fri Apr 28 18:29:07 2000 FLOATING_POINT = IEEE32BIG END HEADER

5.7 Serial I/O

We have spent some time dealing with parallel I/O. It should of course be obvious, that since saving and loading of parallel data is not done from the user program, that the user program must stop its computation when parallel IO needs to be performed. This leads to the model of running Monte Carlo simulations where observables are computed on the fly rather than the one where one just produces configurations, archives them and measures observables later, although this latter model could be uses if driven by a shell script say.

Naturally after computing observables on the fly the user would wish to store them in a file on the front end. Generally (except for propagators and things) observables tend to be small objects that can fit into the memory of a single node.

The QCDSP provides the standard C file access mechanism for reading and writing to files. The C++ iostream class is not implemented. Hence one can use fprintf and fscanf to read from a file and printf and scanf to read from a terminal. (Bob are the binary file functions fread and fwrite implemented?)

There are however several things that the user has to be aware of:

- There is no concept of stdin, stdout and stderr streams on the QCDSP (although there is on the front end). These are UNIX concepts. On the QCDSP one can either write to the screen or to a file.
- The printf command will output on all the processors. The output of processor 0 is sent to the terminal. The output of other processors is buffered and can be retreived using the aprintf command.
- Access to files on the front end is currently available from node 0 only. If other nodes wish to read or write files. They must arrange for node 0 to do it and transfer the results.
- One can turn off I/O by running a program with the command <code>qrun_no_io</code>. In this case I/O produced by <code>printf</code> and <code>fprintf</code> commands is lost.

5.8 Summary

In this section we have discussed the main parallel and serial I/O capabilities of the QCDSP. We now summarise some of the chief results of this chapter.

Memory Summary

User code should reside between addresses 001000 and 05FFFF. The memory layout of a compiled and linked code is given by the appropriate map file. This file also includes the heap area so the highest used address in the map file will be the highest address used by the program as there will be no runtime relocation of symbols, since the linker has already resolved everything.

Parallel I/O General Summary

Parallel I/O has to be performed before execution of user programs or after the completion of user programs. Parallel input involves placing data directly into the QCDSP node memory above the user code but below the operating system. Parallel output proceeds by reading directly the memory of the nodes after a program has stopped executing.

Parallel output can proceed via the gread command for NTFs or via the gunload_lattice command for lattice gauge fields.

Parallel input can proceed via the qload command for NFTs or via the qload_lattice command for lattice gauge fields in the Gauge Connection archive format.

Summary of gread

The basic command is <code>qread <address> <no of blocks></code> which will cause qread to read <no of blocks> data words to be read from each processors memory starting at address <address>. Both <address> and <no of blocks> are hexadecimal numbers without any leading 0x. The behaviour of <code>qread</code> can be controlled by the following QC-Shell Commands

- qset_read_output_filename: Set name of file to dump to. If no arguments are supplied the command prints the currently set name of the dump file. If the command is run with one argument it must be the name of the dump file. Bear in mind that there is a file suffix convention (see figure 5.4.
- qset_read_output_file_access: Set the access mode of the dump file. If called with no arguments the currently set mode is displayed. If called with one argument the argument must be a valid access more. Valid access modes are w for write mode and a for append mode.
- qset_read_output_select: Select output to which the dump will be sent. Dump output can be sent to the screen or to a file. If called with no arguments the command displays the current setting. If called with 1 argument, the argument has to be either f or s for dumping to a file or screen respectively.
- qset_read_output_tagged: Select the tagging mode for the dump if output is going to a file (Node Tagged File). Valid tagging modes are

```
no Do not tag nodes.

yes Use default node tagging mode (tree)

tree Use tree node tagging mode

p4d Use physics (re-mappable) 4D coordinate tagging mode.

m4d Use machine (non-remappable) 4D coordinate tagging mode.
```

Summary of gload

The qload command can be used either to load NTFs or to poke individual words into memory. It can be invoked in three separate modes

- qload <address> <data> -loads <data> to address <address> on all processors.
 Both <data> and <address> are hexadecimal values without any preceding 0x in the
 notation.
- qload <file> will cause the loading of the NTF <file>. The memory addresses are hardwired into the NTF so manual relocation may be necessary.
- qload will cause qload to execute according to the behaviour set with previous commands. Commands affecting qload are
 - qset_load_select Takes 0 or 1 arguments. Selects whether qload operates in screen mode (poking a single data into an address) or file mode (reading an NFT). When called without an argument, the command will print the current selection. The mode will be altered by calling qload with multiple arguments to load files or data directly.
 - qset_load_screen Can be called with 2 or no arguments. If called with two arguments, argument 1 must be a hexadecimal address and argument 2 must be a hexadecimal data value. A subsequent call to qload will load the data value to the address on all processors if screen mode has been selected with qset_load_select, unless the behaviour of qload is altered by some other command.
 - gset_load_filename Can be called with 1 or no arguments. If called with 1 argument, the argument must be a filename. A subsequent call to gload will attempt to load the specified file if gload has been put into file mode using gset_load_select, unless the behaviour of gload has been altered by some other subsequent command.

XS

Node Tag File (NTF) Summary

Node tag files (NTFs) are ASCII files containing data from all the nodes of the currently selected set of processors. A given NTF contains a magic number and a sequence of records, one for each

processor. Each record contains some data to identify its processor, the starting address and the number of blocks of that processors data followed by the data itself. The record is terminated by a 0. Each entry in an NTF needs to be on a new line. All numbers are in hexadecimal. For a summary of Node Tag Files see figures 5.2, 5.3 and 5.4.

Summary of qload_lattice

The command <code>qload_lattice</code> takes two arguments. Argument one is the address on each processor where the first word of lattice data is to reside. The second argument is the name of a lattice gauge configuration file in the Gauge Connection archive format.

Executing the command will cause qload_lattice to distribute the lattice amongst the currently selected nodes. Any conversion from IEEE floating point representation to DSP representation is done at this time. Various diagnostic information is also printed to the screen, including the value of the plaquette and the average of the traces of the gauge field link matrices.

Summary of qunload_lattice

The command qunload_lattice takes no arguments. Details of the lattice have to be made known to it previously using the qset_unload_lattice command.

The qset_unload_lattice command can take 0, 1 or 2 arguments. If qset_unload_lattice is called with no arguments it will display its current settings. If called with 1 argument the argument has to be either d in which case the default settings will be chosen or m in which case the user will be prompted to enter details manually. If called with 2 arguments the first argument has to be an f followed by the name of a file containing a description of the lattice to be dumped. This file must be in the format described in section 5.6.

Gauge Connection Format Summary

See http://gcd.nersc.gov.

Serial I/O Summary

The C Standard I/O file interface is partially implemented. In particular printf, scanf, fprintf, fscanf, fopen, fclose, fread and fwrite are implemented (I am not sure of the last 2). Apart from printf calling these functions only makes sense from node 0. The printf output of processors are saved in a buffer and can later be recovered using qprintf unless the program was run with qrun_no_io. in which case all serial I/O is lost.

Miscellany

In this section we have used rather a lot of the QC-Shell (QOS) built in commands. It is possibly helpful to know that a brief summary of all the QC-Shell commands can be listed using the QC-Shell command qhelp. It is recommended that the output is piped through a pager such as less or more.

Chapter 6

Compiling, Linking and Libraries

6.1 Overview

Until now, we have been able to do all our compilation and linking with the default makefiles. In the last chapter I discussed that for QCDSP applications we may also need an idea of where free memory starts (where I mean free in the sense, that runtime dynamic memory allocation shall not encroach upon it), so that we can use the commands for parallel IO effectively. This information is given by the linker *map* file as discussed previously.

I have mentioned briefly that our control of the stack and the heap sizes is affected through the use of a *linker control file*. This file is also important as it lets us define so called *overlays* which we need to be able to run subroutines in on chip RAM.

What this chapter does not do

However, **this chapter is not intended to be a full Tartan C++ manual**. There is a perfectly good Tartan C++ Manual out there¹ which explains everything in gory detail.

Secondly this chapter is not documentation for the UNIX make system. There are hundreds (thousands) of other documents describing the workings of make available out there².

However we will explain the compiler flags used in the default makefiles, and also mention a few other switches which might be useful.

¹This product is no longer supported by Texas Instruments but it may be "borrowed" from the Columbia University Group

²One of my favourite books explaining all this stuff used to be Graham Glass published by Prentice–Hall. This is now in a second edition. I do not know how good the second edition is

What this chapter does cover

Basically, we would like to cover the following areas in this chapter

- What tools are available
- What do these tools do
- What comprises your program.
- How to control linking via linker control files
- Where you can find exemplar files that you can hack.

6.2 Tartan C++ Software Tools

The Tartan C++ Software suite includes a supposedly industry standard C++ compiler (although its probably not compliant to the level of the latest ANSI standard. Very few C++ compilers are) a linker and a set of object file utilities including an archiver for creating libraries. We now describe some of these utilities in turn.

6.3 The C++ Compiler

The compilation and linking environment is controlled through the so called Tartan C++ Shell. However this is not an interactive shell in the UNIX sense. Rather one should think about it as akin the standard cc command that in turn calls the preprocessor, the various compilation stages and then finally the linker to link all the object files and libraries.

The command name for invoking the Tartan C++ shell is topp. Invoking the command with no arguments lists a help-screen giving all the command line options and flags. According to the manual, the standard invocation is

tcpp options... filename...

where

- options... is zero or more options defining the behaviour of the compilation and linking
- *filename*... is on or more C or C++ source files, assembly source files, object files or library files.

File Suffices

Certain file suffixes are recognised by the shell. For example files with suffices .C or .cpp are treated as C++ files. Files with suffix .c are assumed to be C source files. The full set of standard suffices are listed in figure 6.1.

Extension	File Type	Input or Output
.cpp, .C	C++ Source file (.C recognised under UNIX only)	Input
. C	C Source file	Input
.asm	Assembly source file	Input and output
.tof	TOFF (Tartan) object file	Input and output
.obj	COFF (Texas Instruments) object file	Input
.olb	TOFF (Tartan) library file	Input
.lib	COFF (Texas Instruments) object file	Input
.lcf	Linker Control File	Input
.ctl	Object file list used with linker control file	Output
.i	Preprocessed file	Output
.map	Linker map file	Output

Figure 6.1: The conventional file suffixes for the Tartan Shell. Input or output refers to whether the files are input or output of the shell

Some Common Flags

We now list a few flags commonly used by the shell.

- -c Compile / assemble only. Input source files are compiled into TOFF . toff object files
- -fa file Treat file as an assembly source file
- -fc file Treat file as a C source file
- -fp file Treat file as a C++ source file
- -fl file Treat file as a linker control file
- -e file Produce executable named file. Analogue of -o flag on Standard UNIX compilers.
- -dname =

value

Define preprocessor variable name.

-uname Undefine preprocessor variable name.

- -idir Include directory list dir in the default include path. A directory list is a colon (:) separated list of directories. Analogue of the -I flag on standard UNIX compilers.
- -1k All C files are assumed to be in K&R style.
- -lc All C++ files are assumed to be in cfront mode.
- -ln Normal mode. C and C++ files may contain Tartan's "normal" extensions to the ANSI/ISO standard (default).
- -ka Keep assembly files.
- -kc Keep local control file (.ctl).
- -ke Keep the compiler error file
- -kl Keep the compiler listing and cross reference file (.lst).
- -km Keep the linker map file.
- -v30 Generate code for TMS320C3x DSP processors (defaul). These are the ones we use.
- -1xs I see this in the Makefile but can't find in manual. Bob can you help out?
- -mb Use big memory model.
- -00 Use no optimisation.
- -o1 Use low optimisation.
- -o Optimise; space and time (memory and speed) optimisations are balanced (defaul).
- -os Use space (memory) optimisation.
- -ot Use time (speed) optimisation.
- -q Be quiet; display the minimum amount of information.

Sensible/Necessary Compiler Flags

We need to produce (.ctl) files if we are to link. It is not unhelpful to keep these. We also definitely want a map file at the end of the day so we need to set that. It is often useful to have the listing file (.lst) file around as well and we use the big memory model so a reasonable minimal set of compiler flags is "-km -kc -kl -mb". Hence I usually have the following Makefile macros and rules for compilation.

```
TCPP = tcpp
TCPPFLAGS = -km -kc -kl -mb -q
# Set this to your own include path
INCFLAGS = -i.
.SUFFIXES: .C .tof
.C.tof:
$(TCPP) $(INCFLAGS) $(TCPPFLAGS) -c $<</pre>
```

6.4 QCDSP Runtime and IO Libraries

The QCDSP runtime library is called tcrt30bs.olb and lives in the directory /usr/local/tartan/v2.1 This program is usually automatically linked into your applications so you need not worry about it

The QCDSP I/O libraries (necessary if you want to go and use things like printf) is called tcio30bs.olb. This file is in the default link path but is not automatically linked at run time. You have to do it manually. If you do not link this file you will not be able to do serial IO from your programs. This file is usually linked in by the default makefiles if you use those, but be aware that you have to link it by hand if you start building your own makefiles.

6.5 Converting TOFF to COFF

The Tartan C++ compiler in general produces code in the TOFF object format. However the QCDSP nodes prefer executables to be in the Texas Instruments COFF format. A utility t2c is available to convert from one form to the other. The invocation syntax is:

where the *TOFF* file is a Tartan object file, and the *COFF file* is the desired name of the executable. The filename convention is for executable files to have the .out file suffix.

6.6 A typical Makefile

A typical makefile (named Makefile) to produce an executable would start to look something like this:

```
# Name of makefile so we can recompile if changed
MAKEFILE = Makefile
# The tartan C++ compiler
TCPP = tcpp
```

```
# The minimal set of flags
TCPPFLAGS = -km -kc -kl -mb -q
# My custom include directories
INCFLAGS = -i .:./include
# Base (Unsuffixed) name of the executable
EXEC_BASE = my_executable
# My C++ Sources
SRCS
           = foo.C bar.C main.C
# My Include dependencies
INCLUDES = foo.h bar.h
# My object files: (all my .C files with suffix changed to .tof)
TOF_OBJS = $(SRCS:.C=.tof)
# My Custom libraries
LIBS = mylib1.olb mylib2.olb
# System IO library (Run time library is linked automaticaly)
SYSLIBS = tcio30bs.olb
# Define suffixes
.SUFFIXES: .C .tof
# Build executable
$(EXEC_BASE).out: $(MAKEFILE) $(INCLUDES) $(TOF_OBJS) $(LIBS) $(SYSLIBS)
$(TCPP) -e $(EXEC_BASE).outtof $(INCFLAGS) $(TCPPFLAGS) \
        $(SYSLIB) $(LIBS)
t2c $(EXEC_BASE).outtof $(EXEC_BASE).out
# Default C++ to .tof compilation rule
$(TCPP) $(INCFLAGS) $(TCPPFLAGS) -c $<
# Remove pesky intermediate files
clean:
rm -rf *.tof *.lst *.map core *.outtof *.ctl
```

6.7 Linker Control Files

Linker Control files allow us to have fine tuning over how the linker links our numerous object files into an overall executable. Amongst other things it lets us change the size of the stack and the heap, and allows us to define overlays to use the on chip memory of a node (The so called CRAM).

Before going further a word or two must be said about the linker. It treats your program as a series of modules, containing sections. A module from the linker point of view is an object file. A section can be a defined data section, or a given function entry point. The linker contro file not only gives you power over the memory layout of the program but also, you can individually select (should you so choose) which sections to use, not use etc.

Perhaps the easiest way to explain a linker control file is to examine the default one. You can find it in /usr/local/tartan/v2.1/etc/qcdsp_v5.3.3. The actual file itself is called link.lcf.

I list it below:

```
/*
//
// C and C++ Linker Control File
// for the TI TMS320C30 Application Board
// Startup Configuration
//
```

```
// Modified by RDM 9/8/96 for qcdsp. Assume here that the
// boot kernels have already been used to initialize the hardware,
// including the NGA.
//
// Modified by RDM 10/9/97.
// Changed banner to read version 5
//
*/
link
// Display Banner
banner "Tartan Linker for QCDSP version 5.3.2, modified 1/11/99 by RDM";
banner "Memory = 0x80000, stack = 0x1000, heap = 0x1000";
// Options
resolve sections;
                                     // optimize linking
options romcopy;
                                     // init RAM variables from ROM
// Configuration Constants
define __STACK_SIZE
                    = 0x1000;
                                    // 1K stack
define __SYSMEM_SIZE = 0x1000;
                                     // 1K heap
// These constants define bus control register values loaded at startup
// define tc_EBCR_startup = 0x0;
                                     // expansion bus control value
// These constants define ST, IE registers during initialization
define tc_ST_startup = 0x2800;  // status register
                     = 0x7FFF;
                                    // interrupt mask
define tc_IE_init
// System Configuration
define TDB_BREAKPOINT = 0x66000000;
                                    // breakpoint for Tartan debugger
control "*.ctl";
                                     // list of obj/lib files to link
                                     // *.ctl file built by Tartan shell
list "*.map";
                                     // default extension for link maps
space .stack, __STACK_SIZE;
                                     // declare memory for program stack
space .sysmem, __SYSMEM_SIZE;
                                     // declare memory for program heap
// Explicit Modules
use module = qcdsp_tcroot;
                                     // general startup code
use module = tcrpc;
                                     // needed for Tartan debugger
                                     // link in .cinit tables
use section = .cinit nowarn;
use kind = debug;
                                     // include debug information
                                     // include debug strings
use kind = dstring;
use kind = linenumber;
                                     // include debug line mapping
```

```
// Define Memory Layout
memory (
       limit = 0x80000;
                                       // 0x1000 - 0x7FFFF
       avoid = 0x0, 0x1000;
// Define Program Layout
allocate ".bss" kind = data image (
               origin = 0x1000 kind=data
allocate ".data" kind = constant image (
               kind=constant
allocate ".text" kind = code image (
               kind = code
       );
// Debug Directives
// These commands are for debugger symbol information, and do
// not affect the size or location of the program
allocate debug_directives image kind = debug
        (origin = 0 kind = debug);
allocate debug_strings image kind = dstring
       (origin = 0 kind = dstring);
allocate debug_strings image kind = dstring
       (origin = 0 kind = dstring);
allocate debug_source_location image kind = linenumber
        (origin = 0 kind = linenumber);
end link;
```

You can see that the LCF has a C like command syntax. This file begins with a bunch of comments. The banner instructions are displayed on screen at link time so as to supply you with information about what is going on.

After the banner section you come to the options section. The first command there resolve sections is an instruction to the linker that it should try to optimize the link by missing out (eliminating) sections that have not been referenced. The options romcopy command indicates a particular way of initialising constants. I tend to leave these two well alone.

Next comes the configuration section where we can set the stack and heap sizes. These are done through defining system labels __STACK_SIZE for the stack and __SYSMEM_SIZE for the heap. You can see that the default link file specifies a measly 1K for stack and 1K for heap. In the Columbia physics system __STACK_SIZE is usually set to 0x010000 (64K) and __SYSMEM_SIZE is usually set to 0x030000 (192K). The following symbol definitions deal with breakpoints and interrupts and are probably best left alone.

Then follow two commands control "*.ctl" and list "*.map". The first of these has been built by the compiler during the compilation process. It contains statements like "WITH xxxxxx" where the xxxxx refers to section names. By specifying that the control should come from all the .ctl files, I as a user don't have to type "WITH xxxx manually for each section I want linked into the file. Instead all the referenced sections will be automatically linked for me. The list command merely specifies that the default extention for output map files is .map.

Following these statements come the commands that reserve space for the stack and the heap. These are space .stack, __STACK_SIZE and space .sysmem, __SYSMEM_SIZE. They create two sections named .stack and .sysmem and these will each have the requisite number of words reserved for them.

After this are a couple of use statements. These include entire modules that are I presume part of the run time environment. All this is best left alone.

The command

```
memory (
	limit = 0x80000; //
	avoid = 0x0, 0x1000; // 0x1000 - 0x7FFFF
);
```

defines the virtual address space. This happens to be the same as the physical address space except when you deal with overlays – see next chapter. The statement is that we have a limite of 0x080000 words which sets the upper limit of memory to 0x07FFFF. The second line states that we must avoid locating code between 0x000000 and 0x001000 which is the lower 1K of the memory.

Thereafter we define the layout of the program using allocate statements. Here only the standard sections were allocated, namely the BSS area the constant data area and area for the program text. The total allocation will include the stack and the heap as well. The first allocate command specifies an explicit origin, so that the BSS will start at 0x001000,. The other sections will follow it in memory. Note that each allocated section also has a kind value specifying whether it holds data, constant data or code.

This control file also contains some allocate statements to do with debugging, before it is finished off with the final command END LINK;

Linking With a Linker Control file

To link with a particular linker control file add it to the tcpp command with the -fl flag. A Makefile that uses a linker control file might look something like this:

```
# Name of makefile so we can recompile if changed
MAKEFILE = Makefile
# The tartan C++ compiler
TCPP = tcpp
# The minimal set of flags
TCPPFLAGS = -km -kc -kl -mb -q
```

```
# My custom include directories
INCFLAGS = -i .:./include
# Base (Unsuffixed) name of the executable
EXEC_BASE = my_executable
# My C++ Sources
SRCS
           = foo.C bar.C main.C
# My Include dependencies
INCLUDES
          = foo.h bar.h
# My object files: (all my .C files with suffix changed to .tof)
TOF_OBJS
          = $(SRCS:.C=.tof)
# My Custom libraries
T<sub>1</sub>TBS
      = mylib1.olb mylib2.olb
# System IO library (Run time library is linked automaticaly)
SYSLIBS
           = tcio30bs.olb
# Set up the flags to specify a linker control file
LINK_CONTROL= -fl mylink.lcf
# Define suffixes
.SUFFIXES: .C .tof
# Build executable
$(EXEC_BASE).out: $(MAKEFILE) $(INCLUDES) $(TOF_OBJS) $(LIBS) $(SYSLIBS)
(TCPP) -e (EXEC_BASE).outtof (INCFLAGS) (TCPPFLAGS) \
        $(SYSLIB) $(LIBS) $(LINK_CONTROL)
t2c $(EXEC_BASE).outtof $(EXEC_BASE).out
# Default C++ to .tof compilation rule
$(TCPP) $(INCFLAGS) $(TCPPFLAGS) -c $<
# Remove pesky intermediate files
clean:
rm -rf *.tof *.lst *.map core *.outtof *.ctl
```

6.8 The Tartan Archiver

The Tartan software suite comes with an archiver/librarian called olib, which is the analogue of the UNIX ar command. It can be invoked without any arguments in which case it will present a list of its command line options. Otherwise the syntax is similar to the ar command:

```
olib <options> <library file> <files>
```

For example the command

```
olib -create foo.olb fool.tof foo2.tof foo3.tof
```

will create the library file foo.olb containing the modules in fool.tof fool.tof and fool.tof.

6.9 Summary

In this short chapter we have given some of the more important compiler flags for the tartan C++ compiler and explained meanings of the basic default Linker Control file. We also pointed out how

this file could be changed to change the stack and heap sizes used by a given user program.	

Chapter 7

Single Processor Optimisation I – Using CRAM

7.1 Introduction

The greatest performance gains on the QCDSP can be made through judicious use of the fast on–chip RAM that is built onto the processor and through the circular buffer.

The on-chip RAM, referred to from now on as CRAM is 2Kwords worth of memory that can be immediately accessed by the DSP processor. (It has a memory latency of 0 cycles).

The circular buffer can be thought of as another level of fast memory. It can hold 32Kwords and deliver them to the processor with a latency of a single cycle after some initial setup overhead.

Apart from these features, the DSP itself can overlap one floating point addition with one floating point multiplication. This however requires assembly coding of routines.

One can take advantage of the CRAM and the circular buffer from within C++ at the expense of knowing some magic numbers to set the circular buffer modes. However using the CRAM requires the technique of "overlay programming".

What this chapter covers

In this chapter we wish to explain the process of using the CRAM. We will do this in three stages. Firstly, we will construct a simple SU(3) matrix-vector multiplication routine and time it.

We will then explain how to relocate the C++ routine into the fast CRAM, and time the resulting performance.

Finally, we will give an assembler optimised version of the same routine and locate it into CRAM. At each stage, we should see the number of clock cycles needed to perform the routine decrease¹.

¹We will ignore setup costs, although they are important for the case of a single matrix vector multiply, they will

What this chapter does not cover

We do not cover the following:

• DSP Architecture and Assembly language – see the following:

TMS320C3x User's Guide:

http://www-s.ti.com/sc/psheets/spru031e/spru031e.pdf

TMS320C3x/C4x Assembly Language Tools User's Guide:

http://www-s.ti.com/sc/psheets/spru035c/spru035c.pdf. NOTE this site contains details of the Texas Instruments assembler and linker, not the Tartan Linker described in the last chapter.

TMS320C3x General-purpose applications User Guide:

http://www-s.ti.com/sc/psheets/spru194/spru194.pdf

• Use of the Circular Buffer. See future addenda/re-releases of this document.

7.2 Overlay Programming and the Linker Control File

Overlay programming harks back to the days when computers had a limited amount of memory and kept a lot of their code either in memory that was too slow to access or somewhere else. The premise of the method was that when a piece of code was not actually executed, it could be overwritten by another piece which could then be run. Designing a program so that it split neatly into sets of overlays was a magic artform.

Nowadays of course workstations have effectively unlimited amounts of memory and virtual memory management systems in operating systems take care of paging memory and swapping memory out to disk (a kind of automatic overlaying process).

In embedded systems however overlay programming is still common and in fact for the QCDSP it is necessary. The idea is that the overlays are stored in the slow DRAM, but the linker resolves all the addresses in them as if they were in the CRAM. Before use, a user copies the routines from DRAM into the CRAM and calls them there.

Virtual Addresses and the Linker Control File

In the last chapter we mentioned that the ALLOCATE statements in a linker control file refer to the allocation of *virtual addresses*, but it just so happens that by default the virtual addresses and physical addresses are the same. Now our CRAM is *memory mapped* into the the DSP address space. The lower part of the CRAM contain interrupt and status register areas so effectively from our point of view, the lowest address we can use to access the CRAM is 0x809800.

become negligible when the operation has to be performed for many matrices

Hence routines that use the CRAM must be told that their (virtual) addresses begin at 0x809800 or higher. We can do this with an ALLOCATE statement in the linker control file. The following ALLOCATE statement defines an overlay called mvmult_overlay (for a matrix vector multiply routine) which will take relocatable object code and resolve it as if it began at address 0x809800:

The ALLOCATE statement defines a linker section called mvmult_overlay containing code. The addresses in this linker section will start from 0x809800 (the origin) statement and the linker symbol _mvmult_dest has been defined to contain this address (the define statement). This symbol will be visible from our C++ programs (and the leading underscore is important). The kind statement defines that the section will contain code, and the module command means that the code to be resolved lives in the matvecmult module which in this case means that the object code for this section will be in the matvecmult.tof file or any library containing the modules from the matvecmult.tof file.

So far so good. The address space of the matvecmult module will begin at 0x809800, however we do not want the code to be linked here. We want the actual code section to be in the DRAM. Then, we could copy it to 0x80980 before execution. We can't link it to this address because the loader program (that loads our programs to the QCDSP nodes) can only load to addresses in the DRAM which has an upper limit of 0x07FFFF.

Hence we have to tell the linker control file that our virtual and physical memory addresses are no longer the same and specify the order in which linker sections should be laid out in physical memory. This can be done by the physical statement just before the end of our linker control file

Consider the following physical definition:

```
PHYSICAL (
    origin = 0x1000
    section = ".bss"
    next section = ".data"
    next section = ".text"
    define "_mvmult_begin"
    next section = "mvmult_overlay"
    define "_mvmult_end"
    next remaining
);
```

This tells the linker that real (physical) memory starts at 0x1000 and that in the executable image, the .bss section should start at 0x1000, followed by the .data and .text sections. This

should immediately be followed by the code for the mvmult_overlay section, however we wish to define linker symbols _mvmult_begin to point to the address where the overlay section will start and the _mvmult_end symbol to point to the first address after it ends. Finally, the next remaining statement just tells the linker to put all the remaining sections to follow from the address _mvmult_end.

The idea is that the code will be laid out in physical memory starting at 0x1000. If we want to use the overlay we will have to copy the contents of the DRAM from between the addresses of the _mvmult_begin symbol and the _mvmult_end symbol to the address of the _mvmult_dest symbol. We can then call the subroutine and the code will start executing at the CRAM address because that is what the linker set to be the origin of the routine.

To illustrate how the extra ALLOCATE and PHYSICAL statements fit into the scheme of the usual linker control file, we now show the full file. You should be able to find it in /qcdsp/sfw/qos.5.3.3/examp

```
// OCDSP Modified Default Linker Control File
// Taken from /usr/local/tartan/v2.1/etc/qcdsp_v5.3.3/link.lcf
// and the physics system <phys_root>/mem/include/link_p2v.lcf
// Modifed By BJ
* /
link
// Display Banner
banner "Tartan Linker for QCDSP version 5.3.3, modified 1/11/99 by RDM";
banner "Customised by BJ 4/29/2000";
banner "Memory size: 0x80000 (512 Kwords)";
banner "Stack size: 0x10000 ( 64 Kwords)";
banner "Heap size: 0x30000 (192 Kwords)";
banner "-----";
banner "Remaining : 0x20000 (128 Kwords) for code and gloaded data";
// These are standard
                                   // Eliminate unused sections
resolve sections;
options romcopy;
                                    // init RAM variables from ROM
// Configuration Constants
                   = 0x10000; // 64K stack
= 0x30000; // 192K heap
define __STACK_SIZE
define ___SYSMEM_SIZE
// These constants define bus control register values loaded at startup
// define tc_EBCR_startup = 0x0; // expansion bus control value
// These constants define ST, IE registers during initialization
define tc_ST_startup = 0x2800;  // status register
                                    // interrupt mask
define tc_IE_init
                     = 0x7FFF;
// System Configuration
```

```
define TDB_BREAKPOINT = 0x66000000;
                                       // breakpoint for Tartan debugger
control "*.ctl";
                                       // list of obj/lib files to link
                                       // *.ctl file built by Tartan shell
list "*.map";
                                       // default extension for link maps
space .stack, __STACK_SIZE;
                                       // declare memory for program stack
space .sysmem, __SYSMEM_SIZE;
                                       // declare memory for program heap
// Explicit Modules
                                  // general startup code
use module = qcdsp_tcroot;
use module = tcrpc;
                                       // needed for Tartan debugger
use section = .cinit nowarn;
                                       // link in .cinit tables
use kind = debug;
                                       // include debug information
use kind = dstring;
                                       // include debug strings
use kind = linenumber;
                                       // include debug line mapping
// Define Memory Layout
memory (
       limit = 0x80000;
       avoid = 0x0, 0x1000;
                                       // 0x1000 - 0x7FFFF
);
// Define Program Layout
// Create an overlay so that we can put matvecmult into
// On Chip RAM. 0x80980 is start of Memory mapped image
// of CRAM.
allocate "mvmult_overlay" kind = code image (
       origin = 0x809800 define "_mvmult_dest"
       kind = code module = matvecmult
);
allocate ".bss" kind = data image (
               origin = 0x1000 kind=data
allocate ".data" kind = constant image (
               kind=constant
       );
allocate ".text" kind = code image (
              kind = code
       );
allocate image (
       remaining
);
// Debug Directives
// These commands are for debugger symbol information, and do
// not affect the size or location of the program
```

```
* /
allocate debug_directives image kind = debug
        (origin = 0 kind = debug);
allocate debug_strings image kind = dstring
       (origin = 0 kind = dstring);
allocate debug_source_location image kind = linenumber
        (origin = 0 kind = linenumber);
/* Specify physical memory */
physical (
  origin = 0x1000
  section = ".bss"
  next section = ".data"
  next section = ".text"
  define "_mvmult_begin"
  next section = "mvmult_overlay"
  define "_mvmult_end"
  next remaining
);
end link;
```

Changes to C++ Files

The symbols defined in the linker control files can be made apparent to the C++ source files using the extern "C" statements. For example to see the symbols defined in the linker control file above our C++ program sould contain the lines

```
extern "C" void mvmult_dest();
extern "C" void mvmult_begin();
extern "C" void mvmuld_end();
```

The reason for the extern "C" statements, is to do with the C++ compiler. When you define a function like void foo(float) the C++ compiler encodes the name into some others string such as _foo__XXXXX for the linker symbol where XXXX is some code that signifies that this particular version of the foo takes a single float argument. If you then define a void foo(int) (which you are perfectly allowed to do in C++) the compiler would mangle this into some other compiler symbol, say _foo__YYYY. The extern "C" basically instructs the compiler to not do this encoding. Hence, after the declarations above, if the C++ compiler encounters the mvmult_dest it would not encode it but would straightforwardly translate it to _mvmult_dest which is the correct linker symbol that we defined for it in the linker control file.

Likewise we must declare the actual matVecMult function as extern "C" otherwise the C++ compiler will mangle it to something else and the linker will not find a module called matvecmult. Instead it will find something like _matVecMult__FPfN21, which can lead to linking errors. I do this by defining a header file matvecmult.h and #include-ing it both into the matvecmult.C file where the subroutine is defined (remember it has to be in a separate file to

be a separate module) the and main. C file where the routine is called. The header file looks like this:

```
#ifndef MATVECMULT_H
#define MATVECMULT_H
extern "C" void matVecMult(float *, float *, float *);
#endif
```

Why do the symbols need to be declared as void functions? Well in the linker control file we defined them effectively as entry points to functions. These functions take no arguments in this example, but may in later ones.

Copying the Code to the CRAM

Once the linker control file is set up correctly and the linker symbols are visible we simply need to write a subroutine to copy the matVecMult subroutine into the CRAM. This could be done with a single memcpy instruction, but despite the claims of the Tartan manual, I couldn't manage to find the library where memcpy lives. Consequentially I just wrote a loop to do the memcpy which I show below:

Note here the usage of the external symbols. The addresses of the symbols (without any calling parameters or trailing brackets) are used to get at the memory addresses at which they were defined in the linker control file. The results are cast to unsigned integers.

Let us now go through the process of optimising. We will follow the example above and deal with a single SU(3) matrix times 3-vector multiplication.

7.3 Step One: Straight C++

In this section, we shall write a basic vector and matrix class in file called **matvec.h**, a routine to multiply the two together in **matvecmult.C** and finally a **main.C** file where we time the multiplication routines. There is also a file called **defines.h** where we define a few convenient macros.

I present here the header files. Firstly defines.h

#ifndef DEFINES_H

```
#define DEFINES_H
#define RE 0
#define IM 1
#define N_COMPLEX 2
#define N_COLOUR 3
#endif
secondly matvec.h
// matvec.h
// -----
#ifndef MATVEC_H
#define MATVEC_H
#include "defines.h"
// Quick and dirty Complex 3x3 matrix class
// Supports indexing and can get a
// pointer to the matrix class
// -----
class Matrix {
private:
 float data_[N_COLOUR * N_COLOUR * N_COMPLEX];
public:
 float & operator()(int cmpx, int col, int row)
       return( data_[ cmpx + N_COMPLEX*(col + N_COLOUR*row) ] );
 float* data()
      return &data_[0];
};
// -----
// Quick and dirty complex 3 vector
class Vector {
private:
  float data_[N_COLOUR * N_COMPLEX];
public:
  float & operator()(int cmpx, int row)
       return( data_[ cmpx + N_COMPLEX*row ] );
  float * data()
      return &data_[0];
};
#endif
```

and finally matvecmult.h

```
#ifndef MATVECMULT_H
#define MATVECMULT_H
extern "C" void matVecMult(float *, float *, float *);
#endif
```

The corresponding C++ files are **matvecmult.C**:

```
#include "matvecmult.h"
// Define some indexes into the Matrix
// For Loop Unrolling
// First Row (Row 0) Columnwise
#define RE00 0
#define IM00 1
#define RE10 2
#define IM10
#define RE20 4
#define IM20 5
// Second Row (Row 1) Columnwise
#define RE01 6
#define IM01 7
#define RE11 8
#define IM11 9
#define RE21 10
#define IM21 11
// Third Row (Row 2) Columnwise
#define RE02 12
#define IM02
#define RE12 14
#define IM12 15
#define RE22 16
#define IM22 17
// Define Some indices into the vectors
#define REO 0
#define IMO
               1
#define RE1
               2
#define IM1
               3
#define RE2
#define IM2
void matVecMult(float *y, float *a, float *x)
        // First row
        y[RE0] = (a[RE00]*x[RE0] - a[IM00]*x[IM0])
                +( a[RE10]*x[RE1] - a[IM10]*x[IM1] )
                +( a[RE20]*x[RE2] - a[IM20]*x[IM2] );
        y[IM0] = (a[RE00]*x[IM0] + a[IM00]*x[RE0])
                +( a[RE10]*x[IM1] + a[IM10]*x[RE1] )
                +(a[RE20]*x[IM2] + a[IM20]*x[RE2]);
        // Second Row
        y[RE1] = (a[RE01]*x[RE0] - a[IM01]*x[IM0])
                +( a[RE11]*x[RE1] - a[IM11]*x[IM1] )
+( a[RE21]*x[RE2] - a[IM21]*x[IM2] );
        y[IM1] = (a[RE01]*x[IM0] + a[IM01]*x[RE0])
                +( a[RE11]*x[IM1] + a[IM11]*x[RE1] )
                +( a[RE21]*x[IM2] + a[IM21]*x[RE2] );
```

110

```
y[RE2] = (a[RE02]*x[RE0] - a[IM02]*x[IM0])
                +( a[RE12]*x[RE1] - a[IM12]*x[IM1] )
+( a[RE22]*x[RE2] - a[IM22]*x[IM2] );
        y[IM2] = (a[RE02]*x[IM0] + a[IM02]*x[RE0])
                +( a[RE12]*x[IM1] + a[IM12]*x[RE1] )
                +( a[RE22]*x[IM2] + a[IM22]*x[RE2] );
}
and finally the main program in main.C:
#include <stdio.h>
#include <stdlib.h>
#include "matvec.h"
#include "defines.h"
#include <time.h>
#include "matvecmult.h"
void printTime(unsigned long start, unsigned long end, char *string)
        printf("%s: Start Time %lu\n", string, start);
       printf("%s: End Time %lu\n", string, end);
        printf("%s: Difference %lu cycles\n", string, end - start);
}
int main(int argc, char *argv[])
 Matrix a;
 Vector x;
 Vector y;
 int col, row;
  // Initialize the matrix and the vector
 printf("Entered Program. Setting up Matrix and vector\n");
  for(row = 0; row < N_COLOUR; row++) {</pre>
    for(col = 0; col < N_COLOUR; col++) {</pre>
     a(RE,col,row) = row+1;
      a(IM,col,row) = (float)0;
     x(RE,row) = row+1;
     x(IM,row)
                 = (float)0;
      y(RE,row)
                 = (float)0;
      y(IM,row) = (float)0;
   }
 printf("Calling Mat Vec routine \n");
 // Start timing
 unsigned long start_time, end_time;
 start_time = clock();
  // Call to matrix multiplication routine:
  // y <- A x
 matVecMult(y.data(), a.data(), x.data());
  // Stop timing
 end_time = clock();
 printf("Done. Printing Results\n");
```

// Third Row

111

```
// Print out result
for(row = 0; row < N_COLOUR; row++) {
   printf("[ ( %6f, %6f ) ] \n", y(RE, row), y(IM, row));
}
printTime(start_time, end_time, "matVecMult ");
}</pre>
```

The clock() function can be used to time the performance of the routine. You should be able to find this code in /qcdsp/sfw/qos.5.3.3/example/optimise1 with a suitable makefile and linker control file.

Upon compiling and running I get the following output on the nodes (as usual neglecting the Qdaemon messages)

```
Entered Program. Setting up Matrix and vector
Calling Mat Vec routine
Done. Printing Results
[ ( 6.000000, 0.000000 ) ]
[ ( 12.000000, 0.000000 ) ]
[ ( 18.000000, 0.000000 ) ]
matVecMult : Start Time 113454
matVecMult : End Time 113964
matVecMult : Difference 510 cycles
```

The matrix multiply routine as written, performs 66 floating point operations during execution. Currently this takes 512 ticks of the (is it a 25 or 50MHz clock Bob?) clock.

7.4 Step 2: Relocate C++ to CRAM

Our first optimisation is to relocate the matrix vector multiplication to CRAM. We have described all the steps needed to do this and have shown the complete linker control file in section 7.2. We now show the whole main program:

```
void loadMatVecMultIntoCRAM(void)
  unsigned int length = (unsigned int*)&mvmult_end
                        -(unsigned int*)&mvmult_begin;
  unsigned int * uip_src = (unsigned int *)&mvmult_begin;
  unsigned int * uip_dest = (unsigned int *)&mvmult_dest;
  printf("Copying MatVecMult Into CRAM\n");
  int i;
   for(i=0; i < length; i++) {
        *uip_dest++ = *uip_src++;
  printf("Done\n");
}
void printTime(unsigned long start, unsigned long end, char *string)
        printf("%s: Start Time %lu\n", string, start);
       printf("%s: End Time %lu\n", string, end);
        printf("%s: Difference %lu cycles\n", string, end - start);
int main(int argc, char *argv[])
 Matrix a;
 Vector x;
 Vector y;
 int col, row;
  // Initialize the matrix and the vector
 printf("Entered Program. Setting up Matrix and vector\n");
  for(row = 0; row < N_COLOUR; row++) {</pre>
   for(col = 0; col < N_COLOUR; col++) {</pre>
     a(RE,col,row) = row+1;
      a(IM,col,row) = (float)0;
     x(RE,row) = row+1;
x(IM,row) = (float)0;
      y(RE,row) = (float)0;
      y(IM,row) = (float)0;
   }
  }
 loadMatVecMultIntoCRAM();
 printf("Calling Mat Vec routine \n");
  // Start timing
 unsigned long start_time, end_time;
  start_time = clock();
  // Call to matrix multiplication routine:
  //
  // y <- A x
 matVecMult(y.data(), a.data(), x.data());
  // Stop timing
  end_time = clock();
 printf("Done. Printing Results\n");
```

```
// Print out result
for(row = 0; row < N_COLOUR; row++) {
   printf("[ ( %6f, %6f ) ] \n", y(RE, row), y(IM, row));
}
printTime(start_time, end_time, "matVecMult ");
}</pre>
```

Upon running I find that the program the output is

```
Entered Program. Setting up Matrix and vector Copying MatVecMult Into CRAM

Done
Calling Mat Vec routine
Done. Printing Results
[ (6.000000, 0.000000 ) ]
[ (12.000000, 0.000000 ) ]
[ (18.000000, 0.000000 ) ]
matVecMult : Start Time 240616
matVecMult : End Time 240922
matVecMult : Difference 306 cycles
```

Hence we have managed to make a substantial gain in execution time simply by relocating the C++ code to CRAM.

7.5 Step 3: Assembler optimisation

In this final step, we use a wrapper function that calls an assembler routine to perform the matrix multiplication. Since we are not interested in measuring overheads of relocating memory to CRAM just now, this means moving the timing tests into the wrapper. All the code including makefiles and routines shall eventually be available from /qcdsp/sfw/qos.5.3.3/example/optimise3.

Assembling the Assembler Routine

We have not yet had occasion to use the assembler. The assembler in general use (in the default Makefiles etc) is the Texas Instruments assembler, known as asm30. This produces Texas Instruments (.obj) files. To be able to link these using the tcpp command we must prefix the files with the -fi option to indicate that they are in the Texas instruments COFF object format.

I defined the assembler to be used and the respective flags (having blatantly stolen them from the default makefiles) using the following macros:

```
ASM = asm30
ASMFLAGS = -q - lxs - mb - v30
```

and used the following rules to allow assembly:

and the lines that do the linking look like:

```
ASMSRCS = mtv.asm

ASM_OBJS = $(ASMSRCS:.asm=.obj)

$(OUTBASE).out: $(TOF_OBJS) $(ASM_OBJS) $(INCLUDE) $(MAKEFILE)

$(TCPP) $(TCPPFLAGS) -e $(OUTBASE).outtof $(TOF_OBJS) -fi $(ASM_OBJS)

$(SYSLIB) -f1 $(LCF)

t2c $(OUTBASE).outtof $(OUTBASE).out
```

The assembler routine

The assembler routine is shown below. Most of the contents are explained in the comments. We make the point that the symbol | | indicates that the statements before and after the symbol are to be executed in parallel. This feature is used when loading data from DRAM to CRAM and vice versa, as well as in instructions where an addition and multiplication can be overlapped using the separate addition and multiplication feature of the DSP.

```
***********************
    mtv.asm
    Balint Joo May/2000
    derived from mtv.asm written by RDM
       derived from dirac.asm, written by Dong Chen.
    mtv( (FLOAT *) U, (FLOAT **)x )
******************
Υ
              AR0
    .set
ŢŢ
    .set
              AR2
X
    .set
     .set
     .def _mtv
     .def _cram_vector
* Space to hold 2 vectors in CRAM
_cram_vector
     .space 12
* FUNCTION DEF : _mtv
***********
_mtv:
     PUSH
          FP
    LDI SP,FP
***************
     ; *-FP(1) return_addr
```

```
; *-FP(2)
       ; *-FP(3)
; AR5 __address of cram_vector
       PUSH R4
       PUSH
              R5
       PUSHF
              Rб
       PUSHF
             R7
       PUSH
              AR4
       PUSH
              AR5
       PUSH
              AR6
       PUSH
              AR7
*************
* Load data to CRAM
                                   ; load x into AR4
       T.DT
              *-FP(3), AR4
             80h, AR5
                                   ; get 800000 into AR5
             16, AR5
       LSH
              _cram_vector, AR5
       OR
                                   ; mask in lower bits AR5
                                    ; now holds address of _cram_vector
      LDI
             AR5, AR1
                                   ; put _cram_vector address into AR1
      LDI
              *AR4++, AR0
                                   ; dereference x into ARO
       LDF
              *AR0++, R0
                                   ; get first element *x[0] into R0
      RPTS
                                   ; repeat 5 times
      LDF
              *AR0++, R0
                                   ; get next element *x[i] into R0
                                   ; store *x[i] to _cram_vector[i]
              R0, *AR1++
    || STF
              R0, *AR1++
       STF
                                   ; store last element *x[n-1] in
                                    ; CRAM
* Set pointers for U, Y and X. AR4 is loaded into Y before each
 multiply.
              *-FP(2), U
                                   ; set U to point to matrix
      LDI
                                   ; set X to point to _cram_buffer
      LDI
             AR5, X
      LDI
              6, AR4
                                   ; put 6 into AR4
      ADDI
             AR5, AR4
                                   ; add AR4 to _cram_buffer -> AR4
                                    ; AR4 should point to second vector
                                    ; in CRAM = address of Y
      LDI
              5, IR0
************
* ASM FUNCTION DEF : u dot x
************
u_dot_x:
* X2 = U . X1 (Complex 3x3 matrix dot 3 vector)
* Y : X2
                    return with X2
* U : U - 1
                    return with U + 17
* X : X1
                     return with X1
* R0,R1,R2,R3 are used for parallel instructions
* R0,R1 : temporary storage
* R2 : real part of X2
* R3 : imaginary part of X2
* first row of U
      LDI
             AR4, Y
                                   ; set Y to AR4
                               ;R2 = U[0] * X1[0]
;R3 = U[0] * X1[1]
;P0
                     *X, R2
*++X, R3
*X, R0
       MPYF3
             *U,
       MPYF3
             *U++,
                    *X,
                                   ;R0 = U[1] * X1[1]
       MPYF3 *U,
```

```
MPYF3
                *U++,
                         *-X,
                                  R1
                                          R1 = U[1] * X1[0] | R2 -= R0
    \Pi
       SUBF3
                RO.
                         R2,
                                  R2
        MPYF3
                 *U,
                         *++X,
                                  R0
                                          ;R0 = U[2] * X1[2] | | R3 += R1
       ADDF3
                R1,
                         R3,
                                 R3
        MPYF3
                 *[]++.
                                  R1
                         ++X.
                                          R1 = U[2] * X1[3] | R2 += R0
        ADDF3
                R0,
                         R2,
                                  R2
        MPYF3
                         *X.
                *U,
                                  R0
                                          ;R0 = U[3] * X1[3] | | R3 += R1
    | |
        ADDF3
                R1,
                         R3,
                                  R3
        MPYF3
                 *[]++
                         *-X,
                                 R1
                                          R1 = U[3] * X1[2] | R2 -= R0
        SUBF3
                R0,
                         R2,
                                  R2
        MPYF3
                 *U,
                         *++X
                                  R0
        ADDF3
                         R3,
                                 R3
                                          ;R0 = U[4] * X1[4] | R3 += R1
    R1.
        MPYF3
                 *U++,
                        ++X,
                                  R1
                                          R1 = U[4] * X1[5] | R2 += R0
        ADDF3
                R0,
                         R2,
                                 R 2
        MPYF3
                 *U,
                         *X,
                                  R0
                                          ;R0 = U[5] * X1[5] | R3 += R1
        ADDF3
                R1,
                         R3,
                                 R3
        MPYF3
                         *-X,
                                 R1
                 *U++,
    \prod
        SUBF3
                 R0,
                                  R2
                                          R1 = U[5] * X1[4] | R2 -= R0
* second row of U
        MPYF3
                 *U,
                         *--X(IR0),R0
                                          ;R0 = U[6] * X1[0] | R3 += R1
        ADDF3
                R1,
                         R3,
                                 R3
        LDF
                 *++X,
                         R2
                         *Y++
                                          R2 = X1[1]
                                                              | | X2[0] = R2
    STF
                R2,
        MPYF3
                *U++,
                         R2,
                                 R3
                                          R3 = U[6] * X1[1] | X2[1] = R3
    STF
                R3,
                                          R2 = U[7] * X1[1] | R3 += R1
        MPYF3
                 *[].
                         *X,
                                  R2
        MPYF3
                 *U++,
                         *-X,
                         R0,
                                          R1 = U[7] * X1[0] | R2 = R0-R2
        SUBF3
                R2,
                                  R2
        MPYF3
                 *U,
                         *++X,
                                  R0
                                          ;R0 = U[8] * X1[2] || R3 += R1
        ADDF3
                R1,
                         R3,
                                 R3
        MPYF3
                 *[]++.
                                  R1
                         ++X.
                                          R1 = U[8] * X1[3] | R2 += R0
        ADDF3
                R0,
                         R2,
                                  R2
                *U,
        MPYF3
                         *X.
                                  R0
        ADDF3
                         R3,
                                  R3
                                          R0 = U[9] * X1[3] | R3 += R1
                R1,
        MPYF3
                 *U++,
                         *-X,
                                 R1
                                          R1 = U[9] * X1[2] | R2 -= R0
        SUBF3
                R0,
                         R2,
                                  R2
        MPYF3
                 *U,
                         *++X
                                  R0
        ADDF3
                R1.
                         R3.
                                  R3
                                          ;R0 = U[10] * X1[4] | R3 += R1
        MPYF3
                 *U++,
                         *++X,
                                  R1
                                          R1 = U[10] * X1[5] | R2 += R0
        ADDF3
                R0,
                         R2,
                                 R 2
        MPYF3
                 *U,
                         *X,
                                  R0
                                          ;R0 = U[11] * X1[5] || R3 += R1
        ADDF3
                R1,
                         R3,
                                 R3
        MPYF3
                         *-X,
                                 R1
                 *U++,
        SUBF3
                R0,
                                  R2
                                          ;R1 = U[11] * X1[4] || R2 -= R0
* third row of U
        MPYF3
                         *--X(IR0), R0
    || ADDF3
                         R3,
                                 R3
                                          ;R0 = U[12] * X1[0] || R3 += R1
                R1,
                 *++X,
        LDF
                         R2
                                          R2 = X1[1]
                                                                | | X2[2] = R2
    STF
                R2,
        MPYF3
                 *U++,
                         R2,
                                 R3
                                          R3 = U[12] * X1[1] | X2[3] = R3
    \Pi
       STF
                R3,
                         *Y++
        MPYF3
                 *U,
                         *X,
                                  R2
                                          ;R2 = U[13] * X1[1] || R3 += R1
        MPYF3
                 *U++.
                                  R1
                        *-X,
                                          R1 = U[13] * X1[0] | R2 = R0-R2
        SUBF3
                R2,
                         R0,
                                  R2
        MPYF3
                 *U,
                         *++X
                                  R0
                                          ;R0 = U[14] * X1[2] | R3 += R1
        ADDF3
                R1.
                         R3.
                                  R3
        MPYF3
                 *U++
                         *++X,
                                  R1
        ADDF3
                                          R1 = U[14] * X1[3] | R2 += R0
                R0,
                                  R2
                         R2,
        MPYF3
                 *U,
                         *X,
                                  R0
                                          ;R0 = U[15] * X1[3] || R3 += R1
       ADDF3
                R1,
                         R3,
                                 R3
```

```
MPYF3
            *U++,
                    *-X,
                            R1
 || SUBF3
                                    ;R1 = U[15] * X1[2] || R2 -= R0
            R0,
                            R2
                    R2.
     MPYF3
            *U,
                    *++X,
                            R0
                                    ;R0 = U[16] * X1[4] | R3 += R1
    ADDF3
            R1,
                    R3,
                            R3
    MPYF3
            *U++,
                    *++X,
                            R1
                    R2,
                                    ;R1 = U[16] * X1[5] || R2 += R0
    ADDF3
            R0,
                            R2
            *U,
    MPYF3
                    *X,
                            R0
    ADDF3
                                    ;R0 = U[17] * X1[5] || R3 += R1
            R1,
                    R3,
                            R3
                    *-X,
    MPYF3
            *IJ++ .
                            R1
                    R2,
                                    ;R1 = U[17] * X1[4] || R2 -= R0
    SUBF3
            R0,
                            R2
    MPYF3
            *X++, R0,
                            R0
                                    ; dummy multiply
    STF
            R2.
                                    ;X2[4] = R2
                  R3, R3
                                    ;R3 += R1
     ADDF3
            R1.
                    *Y++
     STF
            R3.
                                    ;X2[5] = R3
Copy result back to DRAM
           AR4. ARO
                                    ; Y in CRAM
    TIDT
            *-FP(3), AR4
    LDI
            *++AR4, AR1
                                    ; Y in DRAM
    LDF
            *AR0++,R0
                                    ; get first element of Y into RO
    RPTS
             4
                                   ; repeat 5 times
            *AR0++, R0
                                   ; get next element of Y into RO
  | STF
            R0, *AR1++
                                   ; store previous element to DRAM
            R0, *AR1++
     STF
                                    ; store last element
     POP
            AR7
                                    ; restore registers and stack frame
     POP
            AR6
     POP
            AR5
            AR4
     POP
            R7
     POPF
    POPF
            R6
     POP
            R5
     POP
            R4
     POP
            FΡ
     RETS
```

By my counting this routine performs the multiplication using 73 floating point instructions, however a lot of these are executed in parallel. Also the routine was originally written to be performing 8 SU(3) multiplies in a similar manner to a 4D finite difference operator. I may not have stripped out all vestiges of this properly. However the routine does seem to work. It is called from **main.C** in a similar manner to the last section. We list the complete **main.C** file below:

```
#include <stdio.h>
#include <stdlib.h>
#include "defines.h"
#include "matvec.h"
#include <time.h>
// -----
// Entry points defined by the linker control file
//
// To all intents and purposes these look like entry
// pointers to functions
// -----
extern "C" void mvmult_end();
extern "C" void mvmult_begin();
extern "C" void mvmult_dest();
// -----
// Entry point of Matrix Vector assembler routine
// -----
extern "C" void mtv(float *u, float **vecs);
```

```
// -----
// Routine for loading assembler routine MTV
// into CRAM
// -----
void loadMatVecMultIntoCRAM(void)
 unsigned int length = (unsigned int*)&mvmult_end
                    -(unsigned int*)&mvmult_begin;
 unsigned int * uip_src = (unsigned int *)&mvmult_begin;
 unsigned int * uip_dest = (unsigned int *)&mvmult_dest;
 printf("Copying MatVecMult Into CRAM\n");
 for(i=0; i < length; i++) {
  *uip_dest++ = *uip_src++;
 printf("Done\n");
void printTime(unsigned long start, unsigned long end, char *string){
      printf("%s: Start Time %lu\n", string, start);
printf("%s: End Time %lu\n", string, end);
      printf("%s: Difference %lu cycles\n", string, end - start);
}
// -----
// Wrapper Routine to Call Assembler routine mtv
//
// Does y = Ax
// -----
void matVecMultMtv(float *y, float *A, float *x)
 // -----
 // mtv expects a (float **)vectors
 // so that vectors[0] = x
 // and vectors[1] = y
 // -----
 float *vector_array[2];
 float **vectors = &vector_array[0];;
 vector_array[0]=x;
 vector_array[1]=y;
 // -----
 // Start Timing
 // -----
 unsigned long start_time, end_time;
 start_time = clock();
 // -----
 // Call assembler
 mtv(A, vectors);
 // -----
 // End Timing
 end_time = clock();
 printTime(start_time,end_time, "MTV Assembler");
int main(int argc, char *argv[])
```

```
int col, row;
Vector x,y;
Matrix A;
// Initialize the matrix and the vector
// -----
printf("Setting up Matrix and vector\n");
for(row = 0; row < N_COLOUR; row++) {</pre>
 for(col = 0; col < N_COLOUR; col++) {</pre>
   A(RE,col,row) = row+1;
   A(IM,col,row) = (float)0;
   x(RE,row) = row+1;

x(IM,row) = (float)0;
   y(RE,row) = (float)0;
   y(IM,row) = (float)0;
// -----
// Load MatVec routine into CRAM
// -----
printf("Loading Matrix Vector Routine into CRAM\n");
loadMatVecMultIntoCRAM();
// -----
// Call to matrix multiplication routine:
// y <- A x
// -----
printf("Calling Mat Vec routine \n");\\
matVecMultMtv(y.data(), A.data(), x.data());
// Print out result
// -----
for(row = 0; row < N_COLOUR; row++) {</pre>
 printf("[ ( %6f, %6f ) ] \n", y(RE, row), y(IM, row));
```

You can find all this code in /qcdsp/sfw/qos5.3.3/example/optimise3. On running I get the following output:

```
Setting up Matrix and vector
Loading Matrix Vector Routine into CRAM
Copying MatVecMult Into CRAM
Done
Calling Mat Vec routine
MTV Assembler: Start Time 303414
MTV Assembler: End Time 303587
MTV Assembler: Difference 173 cycles
[ ( 6.000000, 0.000000 ) ]
[ ( 12.000000, 0.000000 ) ]
[ ( 18.000000, 0.000000 ) ]
```

}

7.6 Summary

In this chapter we have reviewed the basic means of performing single procesor optimisations on the QCDSP. These can be made by making use of the 2Kword on chip RAM and the circular buffer.

The circuler buffer is considered as an advanced topic and is left for future addenda/revisions of this document. We have then detailed the use of the CRAM and presented a three stage process to optimisation:

- 1. Step 1 Write the code as usual
- 2. Step 2 Relocate compute intensive (C++) routine to CRAM
- 3. Step 3 Replace compute intensive routine in CRAM with

assembler routine taking advantage of the ability to execute two instructions per cycle.

Overlay Summary

When defining an overlay, define its origin in the allocate statement of the linker control file to be in the CRAM. It is useful to define a compiler symbol for this origin. Since one is now dealing with virtual addresses that are different from physical addresses, one must include a physical statement at the end of the linker control file to specify the layout of sections. Here it is possible to define symbols pointing to the start and end of the overlay which will facilitate its copying to the CRAM.

Overlay code is module based. A module unit is a .toff or a .coff file. Alternatively an archive .olb file can contain a collection of modules. Hence code for different overlays must reside in different files.

Overlays must be explicitly copied to CRAM before used.

Linker symbols and C++ Files

Linker symbols can be made visible to C++ program files through the use of the extern "C" void symbol() statements, This stops the compiler from encoding symbol. Linker symbols need to be defined with a leading underscore, but this underscore is dropped when dealing with them in C++ files.

Assembly Summary

Assembly coded overlays work the same way as C++ overlays, except that the object files are assembled using the asm30 assembler rather than compiled. The asm30 assembler outputs object

files in the Texas Intsruments COFF format (.obj suffix). Such files should automatically be recognised by tcpp and dealt with appropriately. One may however at one's discretion prepend them with -fi flag which forces tcpp to treat them as COFF object files.